scalafmt: opinionated code formatter for Scala

Ólafur Páll Geirsson

School of Computer and Communication Sciences

Master's Thesis

June 2015

Responsible
Prof. Martin Odersky
EPFL / LAMP

Supervisor
Eugene Burmako
EPFL / LAMP

Abstract

Code formatters bring many benefits to software development, yet they can be tricky to get right. This thesis addresses the problem of developing a code formatter for the Scala programming language. We present scalafmt, an opinionated Scala code formatter that captures many popular Scala idioms and coding styles. Scalafmt has been installed over 6.000 times since its initial release 3 months ago, adopted by several open-source libraries and attracted a vibrant community of external contributors. Our work has been limited to formatting Scala code. Still, this thesis presents algorithms and tooling that we believe can be of interest to the development of code formatters for other programming languages.

Contents

Listings

List of Algorithms

List of Figures

List of Tables

1 Introduction

The main motivation of this study is to bring scalafmt, a new Scala code formatter, to the Scala community. The goal is to capture many popular coding styles so that a wide part of the Scala community can enjoy the benefits that come with automatic code formatting.

Without code formatters, software developers are responsible for manipulating all syntactic trivia in their programs. What is syntactic trivia? Consider the Scala code snippets in listings 1 and 2.

<table>
<tr><td colspan="2">Listing 1: Unformatted code</td><td colspan="2">Listing 2: Formatted code</td></tr>
</table>

```
1  object ScalafmtExample {
2    function(arg1, arg2(arg3(
3      "String literal"),
4      arg4 + arg5))
5  }
6
7
```

```
1  object ScalafmtExample {
2    function(
3      arg1,
4      arg2(arg3("String literal"),
5        arg4 + arg5))
6  }
7
```

Both snippets represent the same program. The only difference lies in their syntactic trivia, that is where spaces and line breaks are used. Although syntactic trivia does not alter the execution of the program, listing 2 is arguably easier to read, understand and maintain for the software developer. The promise of code formatters is to automatically convert any program that may contain style issues, such as in listing 1, into a readable and consistent looking program, such as in listing 2. Code formatting brings several benefits to software development.

Code formatting enables large-scale refactoring. Google used ClangFormat[22], a code formatter, to migrate legacy C++ code to the modern C++11 standard[47]. ClangFormat was used to ensure that the refactored code adhered to Google's strict C++ coding style[15]. Similar migrations can be expected in the near future for the Scala community once new dialects, such as Dotty[37], gain popularity.

Code formatting is valuable in collaborative coding environments. The Scala.js project[41] has over 40 contributors and the Scala.js coding style[10] — which each Scala.js contributor is expected to know by heart — is defined at a whooping 2.600 word count. Each contributed patch is manually verified against the coding style by the project maintainers. This

adds a burden on both contributors and maintainers. Several prominent Scala community members have raised this issue. ENSIME[12] is a popular Scala interaction mode for text editors such as Vim and Emacs. Sam Halliday, an ENSIME maintainer, says "I don't have time to talk about formatting in code reviews. I want the machine to do it so I can focus on the design."[17]. Akka[1] is a popular concurrent and distributed programming library for Scala with over 300 contributors. Viktor Klang, a maintainer of Akka, suggests a better alternative "Code style should not be enforced by review, but by automate rewriting. Evolve the style using PRs against the rewriting config.".[23]. With code formatters, software developers are not burdened by syntactic trivia and can instead direct their full attention to writing correct, maintainable and fast code.

1.1 Contributions

The main contribution presented in this thesis are the following:

- scalafmt, a code formatter for the Scala programming language. At the time of this writing, scalafmt has been installed over 6.000 times in a period of 3 months, adopted by several open source libraries and attracted a vibrant community of external contributors. For details on how to install and use scalafmt, refer to the scalafmt online documentation[13].

- algorithms and data structures to implement line wrapping with a maximum column-width limit — the most challenging part of code formatting. This work is presented in section 3.

- tools to develop and test code formatters. This work is presented in section 4.

The scalafmt formatter itself may only be of direct interest to the Scala community. However, we hope that our experience from developing scalafmt — presented in this thesis — can be help to code formatter developers working with other programming languages.

2 Background

This chapter explains the necessary background to understand Scala and code formatting. More specifically, we motivate why Scala presents an interesting challenge for code formatters. We go into details on Scala's rich syntax and popular idioms that introduced unique challenges to the design of scalafmt. We follow up with a history on code formatters that have been developed over the last 70 years. We will see that although code formatters have a long history, a new tradition of optimization based formatters – which scalafmt proudly joins – started only recently in 2013.

2.1 Scala the programming language

Scala[32] is a general purpose programming language that was first released in 2004. Scala combines features from object-oriented and functional programming paradigms, allowing maximum code reuse and extensibility.

Scala can run on multiple platforms. Most commonly, Scala programs compile to bytecode and run on the JVM. With the releases of Scala.js[10], JavaScript has recently become a popular target platform for Scala developers. Even more recently, the announcement of Scala Native[40] shows that LLVM and may become yet another viable target platform for Scala developers.

Scala is a popular programming language. The Scala Center estimates that more than half a million developers are using Scala[31]. Large organizations such as Goldman Sachs, Twitter, IBM and Verizon rely on Scala code to run business critical applications. The 2015 Stack Overflow Developer Survey shows that Scala is the 6th most loved technology and 4th best paying technology to work with[43]. The popularity of Apache Spark[2], a cluster computing framework for large-scala data processing, has made Scala a language of choice for many developers and scientists working in big data and machine learning.

Scala is a programming language with rich syntax and many idioms. The following chapters discuss in detail several prominent syntactic features and idioms of Scala. Most importantly, we highlight coding patterns that encourage developers to write a single large statement over of multiple small statements. In section 3, we will discover that large statements

```
1  def twice(f: Int => Int) = (x: Int) => f(f(x))
2  twice(_ + 2)(6)  // 10
```

```
1  def twice(f: Function[Int, Int]) =
2    new Function[Int, Int]() { def apply(x: Int) = f.apply(f.apply(x)) }
3  twice(new Function[Int, Int]() { def apply(x: Int) = x + 2 }).apply(6)  // 10
```

introduce a challenge to code formatting.

2.1.1 Higher order functions

Higher order functions (HOFs) are a common concept in functional
programming languages as well as mathematics. HOFs are functions that
can take other functions as arguments as well as return functions as values.
Languages that provide a convenient syntax to manipulate HOFs are said
to make functions first-class citizens.

Functions are first-class citizens in Scala. Consider listing 3. The method
`twice` takes an argument `f`, which is a function from an integer to an
integer. The method returns a new function that will apply `f` twice to an
integer argument. This small example takes advantage of several syntactic
conveniences provided by Scala. For example, in line 2 the argument `_ + 3`
creates a new `Function[Int, Int]` object. The function call `f(x)` is in fact
sugar for the method call `f.apply(x)` on a `Function[Int, Int]` instance.
Listing 4 shows an equivalent program to listing 3, without using syntactic
conveniences. Observe that the body of `twice` was expressed as a single
statement in line 1 of listing 3 but as two independent statements in
listing 4.

2.1.2 Immutability

Functional programming encourages stateless functions which operate on
immutable data structures and objects. An immutable object is an object
that once initialized, cannot be modified. Immutability offers several

Listing 5: Manipulating immutable list

```scala
val input = List(1, 2, 3)
val output = input.map(_ + 1)     // List(2, 3, 4)
                  .filter(_ > 2) // List(3, 4)
```

Listing 6: Manipulating mutable list

```scala
val input = List(1, 2, 3)
val output = mutable.ListBuffer.empty[Int] // mutable list
input.foreach { elem =>
  if (elem + 1 > 2) { // filter
    output += elem + 1 // map
  }
}
output // ListBuffer(3, 4)
```

benefits to software development in areas including concurrency and
testing. Listing 5 shows an example of manipulating an immutable list.
Note that each `map` and `filter` operation creates a new copy of the list
with the modified contents. The original list remains unchanged. Listing 6
show the equivalent operation using a mutable list. Observe that listing 5 is
a single statement while listing 6 contains multiple statements.

2.1.3 SBT build configuration

SBT[39] is an interactive build tool used by many Scala projects. SBT
configuration files are written in `*.sbt` or `*.scala` files using Scala syntax
and semantics. Although SBT configuration files use plain Scala, they
typically use coding patterns which are different from traditional Scala
programs. Listing 7 is an example project definition in SBT. Observe that

Listing 7: SBT project definition

```scala
lazy val core = project
  .settings(allSettings)
  .settings(
    moduleName := "scalafmt-core",
    libraryDependencies ++= Seq(
      "com.lihaoyi" %% "sourcecode" % "0.1.1",
      "org.scalameta" %% "scalameta" % Deps.scalameta))
```

Listing 8: Parsing different Scala dialects with scala.meta

```scala
import scala.meta._
dialects.Sbt0137(
  """lazy val root = project.dependsOn(core)
    lazy val core = project""").parse[Source] // OK
dialects.Sbt0136(
  """lazy val root = project.dependsOn(core)
    lazy val core = project""").parse[Source] // Missing blank line
// Default dialect, regular Scala compilation unit
"""lazy val root = project""".parse[Source] // No top-level statements
```

the project is defined as a single statement and makes extensive use of symbolic infix operators. Due to the nature of build configurations, argument lists to can becomes unwieldy long and a single project statement can span over dozens or even hundreds of lines of code.

2.2   scala.meta

Scala.meta[5] is a metaprogramming toolkit for Scala. Before scala.meta, the state-of-the-art metaprogramming facilities relied on the Scala compiler internals. This had several severe limitations such as too-eager desugaring resulting in loss of syntactic details from the original source code. Scala.meta was designed to overcome these limitations and offer a more robust platform to develop metaprogramming tools for Scala. Several key features of scala.meta have made it an invaluable companion in the development of scalafmt. Most notably among these features are dialect agnostic syntax trees, syntax tree serialization, high-fidelity parsing and algebraically typed tokens.

Scala.meta provides facilities to tokenize and parse a variety of different Scala dialects. One such dialect is SBT configuration files, discussed in section 2.1.3. SBT adds custom support for top-level statements in *.sbt files. An attempt to parse an SBT file with the Scala compiler parser will result in a parse error. Working around this issue requires either depending on the SBT parser or reimplementing its parsing logic. To add insult to injury, top-level statements must be separated by a blank line if you use an SBT version lower than 0.13.6; a restriction that was lifted in SBT 0.13.7. Listing 8 show how to scala.meta makes it trivial to accommodate this zoo of nuances. The result after parsing is a dialect agnostic scala.meta tree structure.

```
1  import scala.meta._
2  > """ object Main extends App { self =>
3        println(s"Hello $self!") // This is a comment
4      }""".parse[Source].get.structure
5  res0: String = """
6  Source(Seq(Defn.Object(Nil, Term.Name("Main"), Template(Nil, Seq(Ctor.Ref.Name("App
        ")), Term.Param(Nil, Term.Name("self"), None, None), Some(Seq(Term.Apply(Term.
        Name("println"), Seq(Term.Interpolate(Term.Name("s"), Seq(Lit("Hello "), Lit
        ("!")), Seq(Term.Name("self"))))))))))
7  """
```

The structure of scala.meta trees can be serialized to strip off all syntactic details. Listing 9 shows how to serialize the tree structure of a simple hello world application. For example, observe that the comment has been stripped away. As we discuss in section 4, this feature was instrumental in testing scalafmt.

Tree node types in scala.meta preserve absolute fidelity with the original source file. This means we can obtain all syntactic details from a tree node such as whether a for comprehension uses parentheses or curly braces as delimiters, whitespace positions, comments and other syntactic trivia. The Scala compiler is infamous for desugaring for-comprehensions into `map/withFilter/flatMap` applications during the parse phase. This made it impossible to implement metaprogramming tasks such as code formatting. High-fidelity parsing in scala.meta has been essential for scalafmt because we can't loose critical syntactic details such as whether for-comprehensions are used over `flatMap`s.

Tokens in scala.meta are strongly typed. Traditional object-oriented libraries treat tokens as a single type with multiple methods such as `isComma/isFor` which returns true if a token instance is a comma or a `for` keyword. However, scala.meta leverages algebraic data types in Scala to represent each different kind of token as a separate type. This feature plays nicely with exhaustivity checking in the Scala pattern matcher and enabled design pattern for the *Router* explained in section 3.3.1.

2.3 Code formatters

Code formatting and pretty printing[1] has a long tradition. In this chapter, we look at a variety of tools and algorithm that have been developed over the last 70 years.

2.3.1 Natural language

The science of displaying aesthetically pleasing text predates as early as 1956[19]. The first efforts involved inserting carriage returns in natural language text. Until that time, writers had been responsible for manually providing carriage returns in their documents before sending them off for printing. The motivation was to "save operating labor and reduce human error". Once type-setting became more commonplace, the methods for breaking lines of text got more sophisticated.

Knuth and Plass developed in 1981 a famous line breaking algorithm[25] for LaTeX, a popular typesetting program among scientific academic circles. LaTeX is the program that was used to generate this very document. The line breaking problem was the same as in the 60s: how to optimally break a paragraph of text into lines so that the right margin is minimized. The primitive approach is to greedily fit as many words on a line as possible. However, such an approach can produce embarrassingly bad output in the worst case. Knuth's algorithm uses dynamic programming to find an optimal layout with regards to a fit function that penalizes empty space on the right margin of the paragraph. This algorithm remains a textbook example of an application of dynamic programming[11, 24].

2.3.2 ALGOL 60

Scowen[42] developed SOAP in 1971, a code formatter for ALGOL 60. The main motivation for SOAP was to make it "easier for a programmer to examine and follow a program" as well as to maintain a consistent coding style. This motivation is still relevant in modern software development.

---

[1] This thesis uses the term *code formatting* over *pretty printing*. According to Hughes[20], pretty printing is a subset of code formatting where the former is only concerned with presenting data structures while the latter is concerned with the harder problem of formatting existing source code — the main topic of this thesis.

```
1   (defun factorial (n)
2     (if (= n 0) 1
3         (* n (factorial (- n 1)))))
```

SOAP did provide a maximum line length limit. However, SOAP would fail execution if the provided line length turned out to be too small. With hardware from 1971, SOAP could format 600 lines of code per minute.

2.3.3 LISP

In 1973, Goldstein[14] explored code formatting algorithms for LISP[28] programs. LISP is a family of programming languages and is famous for it's parenthesized prefix notation. Listing 10 shows a program in LISP to calculate factorial numbers. The simple syntax and extensive use of parentheses as delimiters makes LISP programs an excellent ground to study code formatters.

Goldstein presented a *recursive re-predictor* algorithm in his paper. The recursive re-predictor algorithm runs a top-down traversal on the abstract syntax tree of a LISP program. While visiting each node, the algorithm tries to first obtain a *linear-format*, i.e. fit remaining body on a single line, with a fallback to *standard-format*, i.e. each argument is put on a separate line aligned by the first argument. Goldstein observes that this algorithm is practical despite the fact that its running time is exponential in the worst case. Bill Gosper used the re-predictor algorithm to implement GRINDEF[3], one of the first code formatters for LISP.

Goldstein's contributions extend beyond formatting algorithms. Firstly, in his paper he studies how to format comments. Secondly, he presents several different formatting layouts which can be configured by the users. Both relevant concerns for modern code formatters.

2.3.4 Language agnostic

Derek C. Oppen pioneered the work on language agnostic code formatting in 1980[35]. A language agnostic formatting algorithm can be used for a

variety of programming languages instead of being tied to a single language. Users provide a preprocessor to integrate a particular programming language with the algorithm. Oppen's algorithm runs in $O(n)$ time and uses $O(m)$ memory for an input program of length n and maximum column width m. Besides impressive performance results, Oppen claims that a key feature of the algorithms is its streaming nature. The algorithm prints formatted lines as soon as they are input instead of waiting until the entire input stream has been read. A feature which is typically not a concern for modern code formatters. Moreover, Oppen's algorithm shares a worrying limitation with SOAP: it cannot handle the case when the line length is insufficiently large.

Mark van der Brand presented a library in 1996 that generates a formatter given a context-free grammar[45]. Beyond the usual motivations for developing code formatters, Brand mentions that formatters "relieve documentation writers from typesetting programs by hand". The focus on documentation is reflected by the fact that the generated formatter could produce both ASCII formatted code as well at LATEX markup. Since comments are typically not included in a syntax tree, the presented algorithm has an elaborate scheme to infer the location of comments in the produced output. Like Oppen's algorithm, this library requires the user to plug in a preprocessor to integrate a particular programming language into the Brand's library. Unlike Oppen's algorithm, Brand does not consider line length limits in his algorithm.

John Hughes extended on Oppen's work on language agnostic formatting in term of functional programming techniques[20]. Hughes presented a design of a *pretty-printing* library that leverages combinators with algebraic properties to express formatting layouts. Hughes claims that such a formal approach was invaluable when designing the pretty-printing library, which has seen wise use, including in the Glasgow Haskell compiler. Wadler[46] and Chitil[44] extend on Hughes's and Oppen's work in term of performance and programming techniques. However, this branch of work has been limited to printing data structures and not how to format format existing source code.

2.3.5 gofmt

gofmt[16] is a code formatter for the Go programming language, developed at Google. gofmt was released in the early days of Go in 2009 and is noteworthy for its heavy adoption by the Go programming community. Official Go documentation[6] claims that almost all written Go code — at Google and elsewhere — is formatted with gofmt. Besides formatting, gofmt is used to automatically migrate Go codebases from legacy versions to new source-incompatible releases. However, gofmt supports neither a column limit nor an opinionated setting. Line breaks are preserved in the user's input. For example, listing 11 shows a Go program that uses the same layout as the "unformatted" code (listing 1) in the introduction.

Listing 11: Gofmt example input/output

```go
1  package main
2
3  func main() int {
4    function(arg1, arg2(arg3(
5      "String literal"),
6      arg4+arg5))
7  }
```

The output of running gofmt through listing 11 is identical to the input. This un-opinionated behavior may be considered desirable by many software developers. However, this thesis is only concerned with opinionated code formatting.

2.3.6 Scalariform

Scalariform[38] was released in 2010 and is a widely used code formatter for Scala. Like gofmt, Scalariform does an excellent job of tidying common formatting errors. Moreover, Scalariform supports a variety of configuration options. Scalariform is also impressively fast, it can format large files with over 4.000 lines of code in under 250 milliseconds on a modern laptop. However, Scalariform shares the same limitations with gofmt: it lacks a line length and opinionated setting.

Firstly, the line length setting is necessary to implement many popular coding styles in the Scala community. For example, the Spark[48] and Scala.js[10] coding styles have 100 character and 80 character column

limits, respectively. As we see in other code formatters, adding a line length setting is non-trivial and doing so would require a significant redesign of Scalariform.

Secondly, the lack of an opinionated setting makes it impossible to enforce certain coding styles. For example, the Scala.js coding style enforces *bin-packing*, where arguments should be arranged compactly up to the column length limit and indented by 4 spaces. Listings 12 and 13 shows an example of bin packing enabled and disabled, respectively.

Listing 12: Bin-packing	Listing 13: No bin-packing

```
1  // Column 35                    |
2  class Foo(val x: Int, val y: Int,
3      val z: Int)
4
```

```
1  // Column 35                    |
2  class Foo(val x: Int,
3          val y: Int,
4          val z: Int)
```

Since Scalariform preserves the line breaking decisions from the input, Scalariform has no setting to convert formatted code like in listing 13 to the code in listing 12.

2.3.7   clang-format

Daniel Jasper triggered a new trend in optimization based coded formatters with the release of *ClangFormat*[21] in 2013. ClangFormat is developed at Google and can format C, C++, Java, JavaScript, Objective-C and Protobuf code. An impressive number of languages. Figure 1 shows the architecture of ClangFormat. The main components are the *structural parser* and the *layouter*.

ClangFormat employs a structural parser to split source code into a sequence of *unwrapped-lines*. An unwrapped line is a statement that should fit on a single line if given sufficient line length. A key feature of unwrapped lines is that they should not influence other unwrapped lines. The parser is lenient and parses even syntactically invalid code. The parsed unwrapped lines are passed onto the layouter.

The ClangFormat layouter uses a novel approach to implement line wrapping. Each line break is assigned a penalty according to several rules such as nesting and token type. At each token, the layouter can choose to continue on the same line or break. This forms an acyclic weighted directed
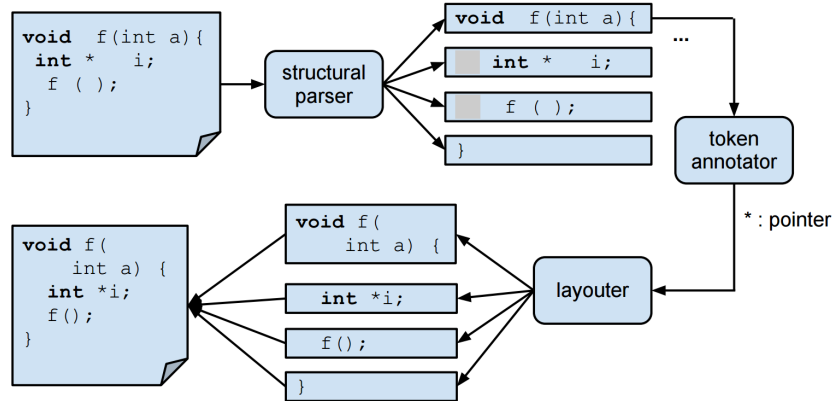
Figure 1: ClangFormat architecture

Listing 14: Unformatted C++ code

```
1  int main(int argc,char const*argv[]) { Defn.Object(Nil, "ClangFormat", Term.Name("
       State"), Foo.Bar( Template(Nil, Seq( Ctor.Ref.Name("ClangLogger")), Term.Param(
       Nil, Name.Anonymous(), None, None)) ), Term.Name("clang-format") ); }
```

graph with tokens representing vertices and splits (e.g., space, no space or
line break) representing edges. The first token of an unwrapped line is the
root of the graph and all paths end at the last token of the unwrapped line.
The layouter uses Dijkstra's[9] shortest path algorithm to find the layout
that has the lowest penalty. To obtain good performance, the layouter uses
several domain specific optimizations to minimize the search space.

Despite supporting an impressive number of programming languages,
ClangFormat does not leverage the language agnostic formatting
techniques described section 2.3.4. Support for each language has been
added as ad-hoc extensions to the ClangFormat parser and layouter.
ClangFormat supports a variety of configuration options, including 6
out-of-the-box styles based on coding styles from Google, LLVM and other
well-known organizations.

A notable feature of ClangFormat is that it's opinionated. ClangFormat
produces well-formatted output for even the most egregiously formatted
input. Listing 14 shows an offensively formatted C++ code snippet.
Listing 15 shows the same snippet after being formatted with
ClangFormat. ClangFormat is opinionated in the sense that it does not
respect the user's line breaking decisions. This feature makes it possible to

```
1  int main(int argc, char const *argv[]) {
2    Defn.Object(Nil, "ClangFormat", Term.Name("State"),
3                Foo.Bar(Template(Nil, Seq(Ctor.Ref.Name("ClangLogger")),
4                                 Term.Param(Nil, Name.Anonymous(), None, None))),
5                Term.Name("clang-format"));
6  }
```

```
1  // Column 35                      |
2  function(
3      firstCall(a, b, c, d, e),
4      secondCall("long argument string"));
```

ensure that all code follows the same style guide, regardless of author.

2.3.8 dartfmt

Dartfmt[29] was released in 2014 and follows the optimization based trend
initiated by ClangFormat. Darftfmt is a code formatter for the Dart
programming language, developed at Google. Like ClangFormat, dartfmt
has a line length setting and is opinionated. Bob Nystrom, the author of
dartfmt, discusses the design of dartfmt in an excellent post[30] on his blog.
In his post, Nystrom argues that the design of a code formatters is
significantly complicated by a column limit setting. The line wrapping
algorithm in dartfmt employs a *best-first search*[36], a minor variant of the
shortest path search in ClangFormat. As with ClangFormat, a range of
domain-specific optimizations were required to make the search scale for
real-world code. Listing 16 shows an example of such an optimization,
avoiding dead ends. The snippets exceeds the 35 character column limit. A
plain best-first search would perform a lot of redundant search inside the
argument list of `firstCall`. However, `firstCall` already fits on a line and
there is no need to explore line breaks inside its argument list. The dartfmt
optimized search is able to eliminate such dead ends and quickly figure out
to break before the `"long argument string"` literal.

2.3.9 rfmt

The most recent addition to the optimization based formatting trend is rfmt[49], a code formatter for the statistical programming environment R. The formatter was released in 2016 – after the background work on this thesis started – and like its forerunners is also developed at Google. rfmt makes an interesting contribution in that it combines the algebraic combinator approach from Hughes[20] and the optimization based approach from LaTeX and ClangFormat.

The algebraic combinator approach makes it easy to express a variety of formatting layouts. rfmt uses 6 layout combinators or *blocks* as they are called in the report. The blocks are the following:

- $TextBlock(txt)$: unbroken string literal.

- $LineBlock(b_1, b_2, \ldots, b_n)$: horizontal combination of blocks.

- $StackBlock(b_1, b_2, \ldots, b_n)$: vertical combination of blocks.

- $ChoiceBlock(b_1, b_2, \ldots, b_n)$: selection of a best block.

- $IndentBlock(n, l)$: indent block b by n spaces.

- $WrapBlock(b_1, b_2, \ldots, b_n)$: Fit as many blocks on each line as possible, break when the column limit is exceeded and align by the first character in b_1.

We'll use an example to show how these relatively few combinators allow an impressive amount of flexibility. Listings 17 and 18 shows two different layouts to format an argument list.

Listing 17: Line block	Listing 18: Stack block

```
1  // Column 35                    |
2  function(argument1, argument2,
3          argument3, argument4,
4          argument5, argument6)
5
```

```
1  // Column 35                    |
2  function(
3    argument1, argument2, argument3,
4    argument4, argument5, argument6
5  )
```

In this case, we prefer the line block from listing 17 since it requires fewer lines. However, our preference changes if the function name is longer as is shown in listings 19 and 20.

Listing 21: Formatting layout for argument lists

```
ChoiceBlock(LineBlock(LineBlock(TextBlock(f), TextBlock("("))),
                      WrapBlock(a1, ... , am),
                      TextBlock(")"),
           StackBlock(LineBlock(TextBlock(f), TextBlock("("))),
                       IndentBlock(4, WrapBlock(a1, ... , am)),
                       TextBlock(")")).
```

Listing 19: Line block

```
1  // Column 35                      |
2  functionNameIsLonger(argument1,
3                       argument2,
4                       argument3,
5                       argument4,
6                       argument5,
7                       argument6)
8
```

Listing 20: Stack block

```
1  // Column 35                      |
2  functionNameIsLonger(
3    argument1, argument2, argument3,
4    argument4, argument5, argument6
5  )
6
7
8
```

Here, we clearly prefer the stack block in listing 20. Listing 21 shows how we use the 6 fundamental blocks in the `rfmt` combinator algebra to express the choice between these two formatting layouts. The variable f denotes the function name and $a_1, ...a_m$ denotes the argument list. Observe that listing 21 does not express how to find the optimal layout.

To find an optimal layout, `rfmt` employs a smart dynamic programming trick. First, it is possible enumerate all layout combinations like the re-predictor algorithm does in section 2.3.3. This leads to exponential growth which turns out to be a problem for some cases. Dynamic programming comes to the rescue by allowing us to reuse partial solutions. Instead of re-calculating the layout cost at each (starting column, block) pair, we store the result in an associative array keyed by the starting column. However, it turns out that this can still be inefficient in terms of memory and speed[2]. To overcome this limitation, Yelland – the `rfmt` author – presents an indexing scheme that makes it possible to extrapolate the layout cost even for missing keys. We refer to the original paper[49] for details. This novel approach enables `rfmt` to format even the most pathologically nested code in near instant time.

---

[2] In fact, ClangFormat started with a similar approach, as explained in this[7] video recording, but then switched to Dijkstra's shortest path algorithms (which in itself is another form of dynamic programming).

Figure 2: Scalafmt architecture

3 Algorithms

This chapter describes how scalafmt formats Scala code. We will see that scalafmt's design is inspired by ClangFormat and dartfmt. However, we believe our design makes a valuable contribution in that it leverages functional programming principles to maximise code reuse and extensibility (TODO, really???).

3.1 Design

Figure 2 shows a broad architectural overview of scalafmt. First, scalafmt parses a source file using scala.meta. Next, we feed a sequence of *FormatToken* data types into a *LineWrapper*. The LineWrapper uses a *Router* to construct a weighted directed graph and run a best-first search to find an optimal formatting layout for the whole file. Finally, the LineWrapper feeds a sequence of *Split* data types into the *FormatWriter*, which constructs a new reformatted source file. The following sections explain these data types and abstractions in detail.

3.2 Data structures

Scalafmt leverages a few carefully designed data structure to allow an implementation that emphasizes correctness and maintainability.

3.2.1 FormatToken

A *FormatToken* is a pair of two non-whitespace tokens. Listing 22 shows the definition of the FormatToken data type.

Listing 22: FormatToken definition

```
1  case class FormatToken(left: Token, right: Token, between: Vector[Whitespace])
```

As shown in the architecture overview in figure 2, each token except the beginning and end of file tokens appear twice in the sequence of FormatTokens: once as the `left` member and once as the `right` member. In a nutshell, the job of the LineWrapper is to convert each FormatToken into a *Split*

3.2.2 Decision

A Decision is a pair of a FormatToken and a sequence of Splits. Listing 23 shows the Definition of decision.

Listing 23: Decision definition

```
1  case class Decision(formatToken: FormatToken, splits: Seq[Split])
```

The *splits* member represents the possible splits that the LineWrapper can choose for *formatToken*.

3.2.3 Policy

A *Policy* is an enforced formatting layout over a region. Listing 24 shows the definition of Policy.

```
1  case class Policy(f: PartialFunction[Decision, Decision], expire: Token)
```

A Policy is a partial function that should be applied to future Decisions up until the *expire* token. Policies easily compose using the Scala standard library `orElse` and `andThen` methods on PartialFunction[3]. Policies enable a high-level way to express arbitrary formatting layouts over a region of code.

3.2.4 Indent

An *Indent* describes indentation over a region of code.

```
1  sealed abstract class Length
2  case class Num(n: Int) extends Length
3  case object StateColumn extends Length
4
5  case class Indent[T <: Length](length: T, expire: Token, inclusive: Boolean)
```

Listing 25 shows the definition of Indent along with the algebraic data type *Length*. Length can either be `Num(n)` where n represents an explicit number of spaces to indent by or `StateColumn` which is a placeholder for the number of spaces required to vertically align by the current column. Indent is type parameterized by Length so that, at some point, we can replace StateColumn placeholders with Nums to obtain a concrete number. For example, given a scala.meta tree `expr`, the definition `Indent(Num(2), expr.tokens.last, inclusive=true)` increases the indentation level by 2 spaces up to and including the last token of `expr`. The `inclusive` member is set to false when the indentation should expire before the expire token, for example in a block wrapped by curly braces, since the closing curly brace should not be indented by 2 spaces. The StateColumn placeholder is required to allow memoization of Splits, which is critical for performance reason as explained in section 3.3.1 on the *Router*.

---

[3] Fun fact. Careful eyes will observe that Policy is in fact a monoid with the empty partial function as identity and function composition as associative operator.

3.2.5 Split

A *Split* represents a (possibly empty) whitespace character to be inserted between two non-whitespace tokens. Listing 26 shows the rather intricate definition of the Split data type[4].

Listing 26: Split definition

```scala
case class Split(modification: Modification,
                 cost: Int,
                 policy: Policy,
                 optimalAt: Option[OptimalToken],
                 indents: Vector[Indent[Length]])(
    implicit val line: sourcecode.Line)
```

The Split data type went through several generations of design before reaching its current structure. Each member serves an important role. The most important member of the Split type is the *modification*. A modification must be one of `NoSplit`, `Space` and `Newline`. The *cost* member represents the penalty for choosing this split. The *optimalToken* member enables an optimization explained in section 3.4.3. The *line* member allows a powerful debugging technique explained in section 3.3.1. The *policy* and *indents* members are explained in sections 3.2.3 and 3.2.4, respectively.

3.2.6 State

A *State* is a partial formatting solution during the best-first search. Listing 27 shows the definition of the State class and companion object.

---

[4] For clarity reasons, a few less important members have been removed from the actual Split definition.

Listing 27: State definition

```
1   case class State(splits: Vector[Split],
2                    totalCost: Int,
3                    policies: Vector[Policy],
4                    indents: Vector[Indent[Num]],
5                    indentation: Int,
6                    column: Int,
7                    formatOff: Boolean) extends Ordered[State] {
8
9     def compare(that: State): Int
10  }
11
12  object State {
13    def nextState(currentState: State, formatToken: FormatToken, split: Split): State
14  }
```

Observe the similarity of State and Split. A State contains various
summaries calculated from the `splits` vector. The summaries are
necessary for performance reasons in the best-first search. Observe that the
indents are type parameterized by `Num`, meaning they only contain concrete
indentations and no `StateColumn` indents. The `indentation` member is the
sum of all currently active indents and `column` represents how many
characters have been consumed since the last newline. The State class
extends the `Ordered` trait to allow for efficient polling from a priority
queue. The `compare` method orders States firstly by their `totalCost`
member, secondly by `splits.length` – how many FormatTokens have been
formatted – and finally breaking ties by the `indentation`. The method
`nextState` calculates a penalty for characters that overflow the column
limit and prepares The method is implemented as efficiently as possible
since the method is on a hot path in the best-first search.

3.3 LineWrapper

The LineWrapper is responsible for turning FormatTokens into Splits. To
accomplish this, the LineWrapper employs a *Router* and a *best-first search*.

3.3.1 Router

The Router's role is to produce a Decision given a FormatToken. Figure 3
shows all possible formatting layout for the small input `val x = y + z`. In
this figure, the Router is the planner that has chosen to open up multiple

Figure 3: Example graph produced by Router

branches at = and + and only one branch for the remaining tokens. This is
no easy task since a FormatToken can be any pair of two tokens. How do
we go about implementing a Router?

The Router is implemented as one large pattern match on a FormatToken.
Listing 28 shows how we can pattern match on a FormatToken and
produce Splits.

```scala
formatToken match {
  case FormatToken(_: Keyword, _) => Seq(Split(Space, 0))
  case FormatToken(_, _: '=')     => Seq(Split(Space, 0))
  case FormatToken(_: '=', _)     => Seq(Split(Space, 0)
                                         Split(Newline, 1))
  // ...
}
```

The pattern _: '=' matches a scala.meta token of type '='. The
underscore _ ignores the underlying value. Keyword is a super-class of all
scala.meta keyword token types. Now, a good observer will notice that this
pattern match can quickly grow unwieldy long once you account for all of
Scala's rich syntax. How does this solution scale? Also, once the match
grows bigger how can we know from which case each Split origins? It turns
out that Scala's pattern matching and scala.meta's algebraically typed
tokens are able to help us.

The Scala compiler can statically detect unreachable code. If we add a case
that is already covered higher up in the pattern match, the Scala compiler
issues a warning. For example, listing 29 shows an example where the
compiler issues a warning.

```
   Listing 29: Unreachable code

1  formatToken match {
2    case FormatToken(_, _: Keyword) => Seq(Split(Space, 0))
3    // ...
4    case FormatToken(_, _: `else`)  => Seq(Newline(, 0)) // Unreachable code!
5  }
```

Here, we accidentally match on a FormatToken with an `else` keyword on
the right which will never match because we have a broader match on a
Keyword higher up. In this small example, the bug may seem obvious but
once the Router grows bigger the compiler becomes unmissable. However,
this still leaves us with the second question of finding the origin of each
Split. Scala macros[4] and implicits[33] come to the rescue.

The source file line number of where a Split is instantiated is automatically
attached on each Split. Remember in listing 26 that the Split case class had
an implicit member of type `sourcecode.Line`. Sourcecode[18] is a tiny
Scala library to extract source code metadata from your programs. The
library leverages Scala macros and implicits to unobtrusively surface useful
information such as line number of call sites. Listing 30 shows how this
works.

```
   Listing 30: Extracting line number from call site

1  Split(Space, 0) /* expands into */ Split(Space, 0)(sourcecode.Line(1))
```

When a `sourcecode.Line` not passed explicitly as an argument to Split's
constructor, the Scala compiler will trigger its implicit search to fill the
missing argument. The `sourcecode.Line` companion contains an implicit
macro that generates a Line instance from an extracted line number. Take
a moment to appreciate how these two advanced features of the Scala
programming language enable a very powerful debugging technique. The
scalafmt router implementation contains 88 cases and spans over 1.000 lines
of code. The ability to trace the origin of each Split has been indispensable
in the development of the Router.

3.3.2 Best-first search

The Decisions from the Router produce a directed weighted graph, as
demonstrated in figure 3. To find the optimal formatting layout, our
challenge is to find the cheapest path from the root node to a final node.

The best-first[36] algorithm is an excellent fit for the task.

Best-first search is a graph search algorithm to efficiently traverse a directed weighted graph. The objective is reach the final token and once we reach there, we terminate the search because we're guaranteed no other solution is better. Algorithm 1 shows a first attempt[5] to adapt a best-first search algorithm to the data structures and terminologies introduced so far. In the best case, the search always chooses the cheapest splits and the

---

Algorithm 1: Scalafmt best-first search, first approach

```scala
/** @returns Splits that produce and optimal formatting layout */
def bestFirstSearch(formatTokens: List[FormatTokens]): List[Split] = {
  val Q = mutable.PriorityQueue(State.init(formatTokens.head))
  while (Q.nonEmpty) {
    val currentState = Q.pop
    if (currentState.formatToken == formatTokens.last) {
      return currentState.splits // reached the final state.
    } else {
      val splits = Router.getSplits(currentState.formatToken)
      splits.foreach { split =>
        Q += State.nextState(currentState, split)
      }
    }
  }
  // Error: No formatting solution found.
  ???
}
```

---

algorithm runs in linear time. Observe that the router is responsible for providing well-behaved splits so that we never hit on the error condition after the while loop. Excellent, does that mean the search is complete? Absolutely not, this implementation contains several serious performance issues.

Algorithm 1 is exponential in the worst case. For example, listing 31 shows a tiny input that triggers the search to explore over 8 million states.

Listing 31: Exponential running time

```
// Column 60                                          |
a + b + c + d + e + f + g + h + i + j + k + l +
m + n + o + p + q + r + s + t + v + w + y +
// This comment exceeds column limit, no matter what path is chosen.
z
```

---

[5] Unfortunately, we make heave use of mutation since graph search algorithms typically don't lend themselves well to functional programming principles.

Even if we could visit 1 state per microsecond — reality is closer to 10 states/microsecond — the search will take almost 1 second to complete. This is unacceptable performance to format only 2 lines of code. Of course, we could special-case long comments, but that would only provide us a temporary solution. Instead, like with ClangFormat and dartfmt, we must apply several domain specific optimizations. In the following section, we discuss the optimizations that have shown to work well for scalafmt.

3.4   Optimizations

This section explains the most important domain-specific optimizations that were required to get acceptable performance for scalafmt. We will see that some optimizations are quite ad-hoc and require some creative workarounds.

3.4.1   dequeueOnNewStatements

Once the search reaches the beginning of a new statement, empty the priority queue. Observe that the formatting layout for each statement is independent from the formatting layout of the previous statement. Consider listing 32.

> **Listing 32: Two independent statements**

```
1  // Column 60                                              |
2  statement1(argument1, argument2, argument3, argument5, argument6)
3  statement2(argument1, argument2, argument3, argument5, argument6)
```

Both statements exceed the column limit, which means that the search must back-track to some extent. However, once the search reaches `statement2` we have already found an optimal formatting layout for `statement1`. When we start backtracking in `statement2`, there is no need to explore alternative formatting layouts for `statement1`. Instead, we can safely empty the search queue once we reach the `statement2` token.

The `dequeueOnNewStatements` optimization is implemented by extending algorithm 1 with an if statement. Algorithm 2 shows a rough sketch of how this is done. With an empty queue, we ensure the search backtracks only as far back as is needed. The `statementStarts` variable contains all tokens

30

Algorithm 2: dequeueOnNewStatements optimization

```
1  // ...
2  val statementStarts: Set[Token]
3  while (Q.nonEmpty) {
4    val currentState = Q.pop
5    if (statementStarts.contains(currentState.formatToken.left)) {
6      Q.dequeueAll // empty search queue
7    }
8    // ...
9  }
```

that begin a new statement. To collect those tokens, we traverse the syntax tree of the input source file and select the first tokens of each statement of a block, each case in a partial function, enumerator in a for comprehension and so forth. The actual implementation is quite elaborate and is left out of this thesis for clarity reasons. Unfortunately, our optimization has one small problem.

Algorithm 2 may dequeue too eagerly inside nested scopes, leading the search to hit the error condition. Listing 33 shows an example where this happens.

Listing 33: Overeager dequeueOnNewStatements

```
1  // Column 50                                          |
2  function1(argument1, { case 'argument2' => 11 }, argument3 // forced newline
3           argument4)
```

Remember that each case of a partial function starts a new statement. The `dequeueOnNewStatements` optimization will dequeue the queue on the first state that reaches the `case` token. In this example, the first state to reach the `case` token will have a strict Policy that disallows newlines up until the closing parenthesis. However, we must insert a newline after the comment. This causes the search to terminate too early and reach the error condition. By inspecting where this problem occurred, we came up with a simple rule to identify regions where the `dequeueOnNewStatements` optimization should be disabled. The simple rule is to never run `dequeueOnNewStatements` inside a pair of parentheses. In section 4, we discuss techniques we used to be confident that this rule indeed works as intended. In the following section (3.4.2) we explain the `recurseOnBlocks` optimization, which allows us to reenable `dequeueOnNewStatements` for selected regions inside parentheses.

3.4.2   recurseOnBlocks

If the `dequeueOnNewStatements` optimization is disabled and we start a new block delimited by curly braces, recursively run the best-first search inside the block. The intuition here is that by recursively running the best-first search, we keep the priority queue small at each layer of recursion. This allows us to run aggressive optimizations such as `dequeueOnNewStatements`.

The `recurseOnBlocks` optimization enables scalafmt to handle idiomatic Scala code where large bodies of higher order functions and blocks are passed around as arguments. Remember from section 2.1 that Scala makes it syntactically convenient to in higher order functions as arguments to other functions. Listing 34 shows an example where this happens and we trigger the `recurseOnBlocks` optimization.

<div>

Listing 34: recurseOnBlocks example

```
1  function(argument1, { higherOrderFunctionArgument =>
2    statement1
3    // ...
4    statementN
5  })
```

</div>

The `dequeueOnNewStatements` optimization is disabled inside argument list. The priority queue grows out bounds because the higher order function can have an arbitrary number of statements.

To implement the `recurseOnBlocks` optimization, we add an extension to algorithm 1. Algorithm 3 shows a rough sketch of how `recurseOnBlocks` is implemented. We change the signature to accept a starting State and token where we stop the search. Observe that we guard against infinite recursion by not making a recursive call on `start.formatToken`. With `recurseOnBlocks` and `dequeueOnNewStatements`, we have solved most problems caused by independent statements affecting the formatting layouts of each other. Next, we leverage recursion again to help the search queue stay small.

Algorithm 3: recurseOnBlocks optimization

```scala
1   def bestFirstSearch(start: State, stop: Token): List[Split] = {
2     val Q = mutable.PriorityQueue(start)
3     while (Q.nonEmpty) {
4       val currentState = Q.pop
5       if (currentState.formatToken.left == stop) {
6         return currentState
7       } else if (currentState.formatToken != start.formatToken &&
8                  currentState.formatToken.left.isInstanceOf['{']) {
9         bestFirstSearch(currentState, closingCurly(currentState.formatToken.left))
10      }
11      // ...
12    }
13  }
```

Listing 35: OptimalToken definition

```scala
1   case class OptimalToken(token: Token, killOnFail: Boolean = false)
```

3.4.3 OptimalToken

An OptimalToken is a hint from a Split to the best-first search that enables the search to early eliminate competing Splits. Recall from listing 26 that a Split has an `optimalToken` member. Listing 35 shows the definition of OptimalToken. When the best-first search encounters a Split with a defined OptimalToken, the best-first search makes an attempt to reach that token with a budget of 0 cost. If successful, the search can eliminate the competing Splits. If unsuccessful and the `killOnFail` member is true, the best-first search eliminates the Split. Otherwise, the best-first search continues as usual.

By eliminating competing branches, we drastically minimize the search space. Listing 36 shows an example where the OptimalToken can be applied. Scalafmt supports 4 different ways to format call-site function applications. This means that there will be 4^N number of open branches when the search reaches `UserObject` N. To overcome this issue, we define an OptimalToken at the closing parenthesis. The best-first search successfully fits the argument list of each `UserObject` on a single line, and eliminates the 3 other competing branches. This makes the search run in linear time as opposed to exponential.

Listing 36: OptimalToken example

```
1  // Column 50                                    |
2  Database(
3      UserObject(name1, age1),
4      UserObject(name2, age2),
5      // ...
6      UserObject(nameN, ageN) // comment will always exceed column limit
7  )
```

To implement the `OptimalToken` optimization, we add an extension to algorithm 3. Algorithm 4 sketches how the extension works. The

Algorithm 4: OptimalToken optimization

```
1   def bestFirstSearch(start: State, stop: Token, maxCost: Int): List[Split] = {
2     // ...
3     val splits = Router.getSplits(currentState.formatToken)
4     var optimalFound = false
5     splits.withFilter(_.cost < maxCost).foreach { split =>
6       val nextState = State.nextState(currentState, split)
7       split.optimalToken match {
8         case Some(OptimalToken(expire, killOnFail)) =>
9           val nextNextState = bestFirstSearch(nextState, expire, maxCost = 0)
10          if (nextNextState.expire == expire) {
11            optimalFound = true
12            Q += nextNextState
13          } else if (!killOnFail) {
14            Q += nextState
15          }
16        case _ if !optimalFound =>
17          Q += nextState
18      }
19    }
20    // ...
21  }
```

`bestFirstSearch` method has a new `maxCost` parameter, which is the highest cost that a new splits can have. Next, if a Split has defined an `OptimalToken` we make an attempt to format up to that token. If successful, we update `optimalFound` variable to eliminate other Splits from being added to the queue. If unsuccessful and `killOnFail` is true, we eliminate the Split that defined the OptimalToken. A straightforward extension to this algorithm would be to add a `maxCost` member to the `OptimalToken` definition from listing 35. However, this has not been necessary for scalafmt.

3.4.4 pruneSlowStates

The pruneSlowStates is a optimization that eliminates states that progress slowly. A state progresses slowly if it visits a token after other equally or less expensive states. The insight is that if two equally expensive states visit the same token, the first state to visits that token typically produces a better formatting layout.

By eliminating slow states, we obtain a better formatting output in addition to minimizing the search space. Listing 37 shows two equally expensive formatting solutions where one solution is fast and other is slow.

Listing 37: Slow states

```
1  // Column 30                    |
2
3  // Fast state
4  a + b + c + d + e + f + g +
5  h + i + j
6  // slow state
7  a + b + c +
8  d + e + f + g + h + i + j
```

Of course, the line break after `g +` could be more expensive than the line break after `c +`. Instead, the Router is free to assign an identical cost to both line breaks and let `pruneSlowStates` take care of eliminating the slow state.

The `pruneSlowStates` is implemented as a extension to algorithm 4. Algorithm 5 shows a rough sketch of how the extension works.

Algorithm 5: pruneSlowStates optimization

```
1   // ...
2   val fastStates: mutable.Map[FormatToken, State]
3   while (Q.nonEmpty) {
4     val currentState = Q.pop
5     if (fastStates.get(currentState.formatToken)
6                 .exists(_.cost <= currentState.state) {
7       // do nothing, eliminate currentState because it's slow.
8     } else {
9       if (!fastStates.contains(currentState.formatToken)) {
10        // currentState is the fastest state to reach this token.
11        fastStates.update(currentState.formatToken, currentState)
12      }
13      // continue with algorithm
14    }
15  }
```

Observe that this algorithm is transparent to the Router. No special annotations are required from Splits.

3.4.5   escapeInPathologicalCases

Alas, despite our best efforts to keep the search space small, some inputs can still trigger exponential running times. The `escapeInPathologicalCases` optimization is our last resort to handle such pathological inputs. How do we detect that the search has encountered such a challenging input?

We detect the search space is growing out of bounds by tallying the number of visits per token. If we visit the same token N times, we can estimate the current branching factor to be around $log_2(N)$. In scalafmt, we tune N to be 256 so that the best-first search can split into two or more paths for up to 8 tokens. When a token has been visited more than 256 times, we trigger the `escapeInPathologicalCases` optimization. In the following paragraphs section, we present two alternative fallback strategies: *leave unformatted* and *best-effort*.

The simplest and most obvious fallback strategy is to leave the pathologically nested code unformatted. This can be implemented by backtracing to the first token of the current statement and then reproduce the formatting input up to the last token of said statement. This method is guaranteed to run linearly to the size of the input. The responsibility is left to the software developer to a manually format her code, removing all the benefits of code formatting. However, in some cases the software developer may prefer to get a decent yet suboptimal formatting output.

The best-effort fallback strategy applies heuristics to give a decent but suboptimal formatting output. When a token is visited for the 256th time, we select two candidate states from the search queue and eliminate all other states. The first state is the fastest state — the state that has reached furthest into the token stream — that is not bound a prohibitive single line policy. Single line policies are policies that eliminate newline Splits. The second state is the current state — the slow state that visited the token for the 256th time. The intuition is that the fast state has good formatting output so far but for is stuck on a challenging token for some reason. The slow may paid a hefty penalty early causing it to move slowly but maybe the early penalty will yield a better output in the end.

Algorithm 6 shows an example of the best-effort strategy can be implemented as an extension to algorithm 1. The `isSafe` method on State

---

Algorithm 6: best-effort fallback strategy

```
1   var fastestState: State
2   val visits: mutable.Map[FormatToken, Int].withDefaultValue(0)
3   while (Q.nonEmpty) {
4     val currentState = Q.pop
5     visits.update(currentState.formatToken, 1 + visits(currentState.formatToken))
6     if (currentState.length > fastestState.length && currentState.isSafe) {
7       fastestState = currentState
8     }
9     if (visits(currentState.formatToken) == MAX_VISITS_PER_TOKEN) {
10      Q.dequeueAll
11      Q += fastestState
12      Q += currentState
13      visits.clear()
14    } else {
15      // continue with algorithm
16    }
17  }
```

---

returns true if the state contains prohibitive policies, derived from annotated metadata in Splits from the Router. Observe that this algorithm will reapply the best-effort fallback until the search reaches the final token. In scalafmt, we bound the number of this can happen with a final fallback to the unformatted strategy.

The unformatted and best-effort fallback strategies offer different trade-offs. The unformatted strategy works well in a scenario where a software developer is available to manually fix formatting errors. The best-effort strategy works well on computer generated code where just a modicum of formatting still greatly aid the legibility of the code. Unfortunately, as we discuss in section 4.2, we struggled to guarantee idempotency using the best-effort strategy. This limitation renders the best-effort strategy useless in environments where code formatters are used to enforce a consistent coding style across a codebase.

3.5   FormatWriter

Recall from figure 2, the FormatWriter receives splits from the best-first search and produces the final output presented to the user. In addition to reifying Splits, the FormatWriter runs three post-processing steps:

docstring formatting, vertical alignment and *stripMargin alignment.*

3.5.1 Docstring formatting

Docstrings are used by software developers to document a specific part of code. Like in Java, docstrings in Scala start with the `/**` pragma and end with `*/`. However, unlike in Java, the Scala community is split on whether new lines inside docstrings should align by the first or the second asterisk. The official Scala Style Guide[8] dictates that new lines should align by the second asterisk while the Java tradition is to align by the first asterisk. The Scala.js[10] and Spark[48] style guides follow the Java convention. To accommodate all needs, scalafmt allows the user to choose either style. To enforce that the asterisks are aligned according to the user's preferences, the FormatWriter rewrites docstring tokens. This is implemented with simple regular expressions.

3.5.2 stripMargin alignment

The Scala standard library adds a `stripMargin` extension method on strings. The method helps Scala developers write multiline interpolated and regular string literals. Listing 38 shows an example usage of the `stripMargin` method.

Listing 38: stripMargin example

```scala
object StripMarginExample {
  """Multiline string are delimited by triple quotes in Scala.
    |You can write as many lines as you want.""".stripMargin
}
```

After calling the method, the indentation and | character on line 3 are conveniently removed. However, the hard-fought indentation on the pipe can easily be lost when the string is moved up or down a scope during refactoring. Scalafmt can automatically fix this issue. In the FormatWriter, scalafmt rewrites string literals to automatically align the | characters with the opening triple quotes `"""`. This setting is disabled by default since scalafmt requires semantic information to confidently determine if the `stripMargin` invocation calls the standard library method or a user-defined method.

3.5.3 Vertical alignment

It turns out that vertical alignment is incredibly popular in the Scala programming community. When lines are vertically aligned when similar operators from a sequence of lines align together. Listing 39 shows an example of vertical alignment.

```scala
object VerticalAlignment {
  x match {
    // Align by => and -> and //
    case 1  => 1  -> 2  // first
    case 11 => 11 -> 22 // second

    // Blank lines separate alignment blocks.
    case ignoreMe => 111 -> 222
  }

  def name   = column[String]("name")
  def status = column[Int]("status")
  val x  = 1
  val xx = 22

  libraryDependencies ++= Seq(
    "org.scala-lang" % "scala-compiler" % scalaVersion.value,
    "com.lihaoyi"    %% "sourcecode"    % "0.1.1"
  )
}
```

Listing 39: Vertical alignment example

The lack of vertical alignment in the initial release of scalafmt was a hindrance for user-adoption. Configurable vertical alignment was added to the 0.2 release.

Vertical alignment is implemented in the FormatWriter as an extension to the reification of Splits. Instead of reifying a Space into a single space literal, the FormatWriter builds a `tokenAligns: Map[Split, Int]` which specifies the number of additional spaces to add for each Space. Listing 7 shows a simplification of how the actual algorithm[6] constructs the `tokenAligns` map. In a nutshell, the algorithm builds blocks of lines that can be vertical aligned. Blocks are separated by blank lines or mismatching candidates in subsequent lines. As demonstrated in listing 39, key features of this algorithm include that users can configure arbitrary symbols to align by and lines can contain multiple vertically aligned columns. This algorithm

---

[6] The actual implementation is 130 lines of code including helper methods, see https://git.io/voIG4

Algorithm 7: Vertical alignment, simplified algorithm

```scala
case class FormatLocation(formatToken: FormatToken, split: Split, state: State)
/** Returns true if location is eligible for vertical alignment */
def isCandidate(location: FormatLocation): Boolean
/** Returns true if all vertical alignment candidates in a and b match */
def allColumnsMatch(a: Array[FormatLocation], b: Array[FormatLocation]): Boolean
/** Returns map where the keys are (0 to block.length) and values are the
    corresponding column index where all candidates should align */
def getMaxColumns(block: Vector[Array[FormatLocation]]): Map[Int, Int]

def getAlignTokens(
    locations: Array[FormatLocation],
    alignConfiguration: Map[String, Regex]): Map[Split, Int] = {
  val finalResult = Map.newBuilder[Split, Int]
  val lines: Array[Array[FormatLocation]] = getLines(locations)
  var block = Vector.empty[Array[FormatLocation]]
  for (formatLocations <- lines) {
    val candidates: Array[FormatLocation] = formatLocations.filter(isCandidate)
    if (block.isEmpty) { // Starting a new block.
      if (candidates.nonEmpty) block = block :+ candidates
    } else {
      if (columnsMatch(block.last, candidates)) {
        block = block :+ candidates
      } else { // release alignment
        val maxColumns = getMaxColumns(block)
        for (line <- block) {
          for ((tokenToAlign, columnIndex) <- line.zipWithIndex) {
            finalResult += (tokenToAlign.split,
                            maxColumns(columnIndex) - tokenToAlign.state.column)
          }
        }
      }
    }
  }
}
```

has two main limitations. Firstly, all lines in a block must have an equal number of matching columns. Secondly, the algorithm can't infer vertical alignment forcing the user to explicitly configure which tokens should align.
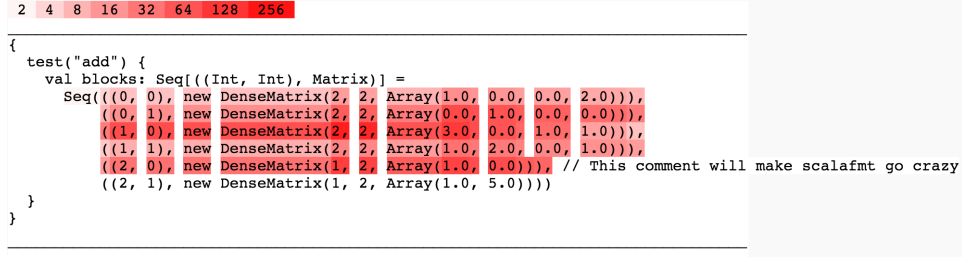
```
2  4  8  16  32  64  128  256

{
  test("add") {
    val blocks: Seq[((Int, Int), Matrix)] =
      Seq(((0, 0), new DenseMatrix(2, 2, Array(1.0, 0.0, 0.0, 2.0))),
          ((0, 1), new DenseMatrix(2, 2, Array(0.0, 1.0, 0.0, 0.0))),
          ((1, 0), new DenseMatrix(2, 2, Array(3.0, 0.0, 1.0, 1.0))),
          ((1, 1), new DenseMatrix(2, 2, Array(1.0, 2.0, 0.0, 1.0))),
          ((2, 0), new DenseMatrix(1, 2, Array(1.0, 0.0))),  // This comment will make scalafmt go crazy
          ((2, 1), new DenseMatrix(1, 2, Array(1.0, 5.0))))
  }
}
```

Figure 4: Example heatmap with 5.121 visisted states

4 Tooling

This chapter describes the tools that we developed while designing an implementing algorithms for scalafmt. These tools were indispensable in giving us confidence that our algorithms worked as intended.

4.1 Heatmaps

Section 3.4 introduces several extensions to algorithm 1 that were required go get acceptable performance. In general, the extensions involved eliminating search states. To identify code patterns that triggered excessive search growth, we developed heatmaps.

Heatmaps are a visualization that displays which code regions are most frequently visited in the best-first search. Figure 4 shows an example heatmap. The intensity of the red color indicates how often a particular token was visited. A token highlighted by the lightest shade of red was visited twice while a token highlighted by the darkest shade of red was visited over 256 times. This figure demonstrates several of the optimizations discussed in section 3.4. Firstly, thanks to the `dequeueOnNewStatements` optimization, the background is plain white up to the `Seq`. The `Seq` gets visited twice, once when there's a space after the = and once when there's a newline. Secondly, due to the OptimalToken optimization, when the search gets into trouble it backtracks to the tuple `(0, 0)` instead of the `Seq[((Int, Int), Matri)]` type signature. Finally, because of the strategically placed comment at the end that exceeds the column limit, the search space grows out of bounds on the fourth argument triggering the `escapeInPathologicalCases` best-effort fallback. Without heatmaps, it would be a much greater challenge to get these insights.

```
2  4  8  16  32  64

List(Split(Space,
           0,
           policy = SingleLineBlock(close),
           ignoreIf = blockSize > style.maxColumn),
      Split(nl, 1, policy = {
           case Decision(t@FormatToken(_, `close`, _), s) =>
             Decision(t, List(Split(Newline, 0)))
         })
```

Figure 5: Example diff heatmap

However, these heatmaps gave us limited insights in how our changes affected the search space in the best-first search.

We developed an extension to heatmaps that allows us to visually compare the difference in search space between two versions of scalafmt. Figure 5 shows an example of such a diff report. The green background indicates that the new version of scalafmt makes fewer visits to those regions. Observe that the > operator has a background with a light shade of red. This means that the operator was visisted more often in the new scalafmt version. A price well worth paying considering the overall shrink in search space. To produce diff heatmaps, we first persist to a database the statistics report needed to generate a single heatmap after each test run. Then, we generate the diff heatmap by fetching two reports and calculating the difference in visits per token. If the difference is negative for a particular token — meaning we visited said token fewer times — the background is highlighted green, otherwise red.

4.2   Property based testing

 Property based tests verify that software implements a high-level specification. Unlike example based tests, the input is not manually curated by the software developer. Instead, the input is automatically generated. Property based tests played a vital role in development of scalafmt and gave us confidence that the algorithms in section 3 behave well in the real world. We collected a sample of 1.2 million lines of code from open source Scala projects available online. The sample was compressed into a 23mb zip

file[7]. Our test suite would download the sample and test again three properties: *can-format*, *AST integrity* and *idempotency*.

4.2.1 Can-format

The can-format property simply says that if the Scala compiler's parser is able to parse the source input file, then scalafmt should be able to format the source file. Although this may seem like a trivial property, it was by far the most effective property at finding bugs in scalafmt. Most commonly, comments in the most unexpected placed would cause the best-first search to not reach the last token in the input. An overly strict Policy was usually the culprit of such bugs, which was easy to fix thanks to our tracing techniques described in section 3.3.1.

4.2.2 AST integrity

The AST integrity property says that the abstract syntax tree of the formatted source file should be identical to the abstract syntax tree of the original input. Recall from section 2.2 that scala.meta trees can be serialized into strings. We leverage this feature to test AST integrity. Algorithm 8 shows the code needed to test AST integrity. This property

---

Algorithm 8: AST integrity property

```
1   import scala.meta._
2   forAll { (code: String) =>
3     val beforeAST = code.parse[Source].show[Structure]
4     val afterAST = Scalafmt.format(code).parse[Source].show[Structure]
5     beforeAST == afterAST
6   }
```

---

catched several critical bugs. For example, in one case, scalafmt inserted a newline after the keyword `return`, breaking the semantics of the original source code. Moreover, this property highlighted the danger of enabling the `stripMargin` alignment. Since the `stripMargin` modified the contents of regular and interpolated string literals, the AST of the formatted output changed. A sign that `stripMargin` should be disabled by default. Finally,

---

[7]See https://github.com/olafurpg/scalafmt/releases/download/v0.1.4/repos.tar.gz

by adhereing to AST integrity, we confine the scope of what scalafmt can do. This makes it easy for us to reject feature requests such as sorting imports. Such refactoring can be left to a separate tool, in our humble opinion.

4.2.3   Idempotency

The idempotency property says that if the output of formatting a source file twice should be identical to the output of formatting it once. This property is critical for scalafmt to be used as part of any continuous integration setup. It is not at all obvious that the algorithms in section 3 fulfill the idempotency property. Our experience reveals that it is in fact very easy to accidentally introduce non-idempotent formatting rules in the Router. We did not test for idempotency until the 0.2.3 release, after users reported non-idempotent formatting behavior in scalafmt. Yet, even after we started testing against idempotency in our comprehensive test-suite, we continued to receive issues with non-idempotent formatting. It turns out that 1.2 million lines of code is not a large enough sample to catch all property bugs.

5 Evaluation

Code formatting is inherently a subjective topic. This introduces a challenge when evaluating a code formatter. In this chapter, we will present measurements that we believe show the success of scalafmt. We do not measure how well software developers perceive scalafmt formatted code. Instead, we will focus on *performance benchmarks* and *user adoption.*

5.1 Performance benchmarks

In this chapter, we describe our test methodology, key metrics and the performance results.

5.1.1 Setup

The benchmarks are run on a Macbook Pro (Retina, 15-inch, Mid 2014) laptop with a quad-core 2.5 GHz Intel Core i7 processor, 256 KB L2 cache per core and 6 MB shared L3 cache. The laptop has 16 GB 1600 MHz DDR3 memory. The operating system is OS X El Capitan 10.11.5. We run the benchmarks from the scalafmt commit id aff5e794 compiled against Scala 2.11.7, running on JVM version 8, update 91. All benchmarks are run with the OpenJDK Java Microbenchmark Harness (JMH)[34]. Benchmarking on the JVM is notoriously hard, JMH is designed to take into account a variety of parameters that affect performance on the JVM to obtain accurate measurements. The sbt-jmh[27] plugin makes it easy to integrate JMH with a Scala project.

5.1.2 Macro benchmark

The macro benchmarks is designed to get insight on how scalafmt performs in a continuous integration setup. We format the entire Scala.js codebase. The codebase contains ??? source files with combined ??? lines of code. We believe the Scala.js codebase strikes a fine balance between size and X. For accurate measurements, we run five iterations of the macro benchmark. We compare the running time with Scalariform. The macro benchmark is multi-threaded and takes advantage of all cores on the underlying hardware.

Benchmark	Score	Error	Units
Parallel.scalafmt	14.616	± 0.632	s/op
Parallel.scalariform	2.810	± 0.641	s/op
Ratio	5.20		
Synchronous.scalafmt	35.654	± 0.459	s/op
Synchronous.scalariform	5.951	± 0.135	s/op
Ratio	5.99		

Table 1: Results from macro benchmark.

25th	Median	Mean	75th	90th	95th	99th	Max
16	49	115	127	275	443	979	4670

Table 2: Percentiles of lines of code per file in macro benchmark.

Table 1 shows the results from the macro benchmark. Scalafmt is much slower, which may not be surprising considering since scalafmt's primary goal has always been to produce appealing output. Scalafmt may try thousands of different formatting layouts to produce an optimal formatting output. Scalariform's formatting algorithm is linear. Still, we are quite content with the performance of scalafmt. This benchmark reveals that scalafmt can format 12.000 lines of code per second.

5.1.3 Micro benchmark

The micro benchmark is designed to get insight into how scalafmt performs in an interactive software developer workflow. For example, it is common to configure a text editor to format source code on every save. First, we find out how many lines of code are typical in a source file.

We performed a small study to learn the size of a typical source file. We collected a sample of 3.2 million lines of code. Table 2 shows the distribution of file sizes in our sample. For example, 25th percentile of files have 16 lines of code. Observe that over 90 percent of all files are rather small, or under 300 lines of code. Less than one percent of all files contain over 1.000 lines of code.

Using the results from our small study, we choose to run the micro

Benchmark	Score	Error	Units
ExtraLarge.scalafmt	1423.140	± 103.360	ms/op
ExtraLarge.scalariform	219.820	± 14.450	ms/op
Ratio	6.50		
Large.scalafmt	355.819	± 17.385	ms/op
Large.scalariform	39.324	± 3.395	ms/op
Ratio	9.05		
Medium.scalafmt	79.616	± 2.013	ms/op
Medium.scalariform	15.934	± 0.441	ms/op
Ratio	5.00		
Small.scalafmt	6.968	± 0.104	ms/op
Small.scalariform	1.176	± 0.025	ms/op
Ratio	5.93		

Table 3: Results from micro benchmark.

benchmark on select four selected files of varying sizes: small ($\sim$ 50 LOC), medium ($\sim$300 LOC), large ($\sim$1.000 LOC) and extra large ($\sim$4.500 LOC). To minimize error margins, we run 10 warmup iterations followed by 10 measured iterations. As in the macro benchmark, we compare the running time with Scalariform. The micro benchmark is single threaded.

Table 3 shows the results from the micro benchmark. No surprise, scalafmt is a slower than Scalariform again. In fact, scalafmt seems a lot slower in this benchmark compared to the macro benchmark. Why is that? Moreover, is this performance gap acceptable for interactive software development.

To answer the first question, we believe scalafmt performs better in the macro benchmark — relative to Scalariform — because macro benchmark is multi-threaded while the micro benchmark is single threaded. As the measurements show, scalafmt operations are slower. Hence, the scheduler in the macro benchmark is presumably more efficient at allocating CPU wall time to the scalafmt test runs compared to Scalariform.

Secondly, we believe this performance is usable for most use-cases, but needs improvements to be acceptable for settings such as reformat on save/compile. We currently don't recommend users to configure text
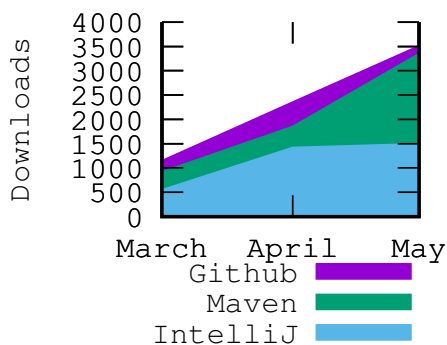
Figure 6: Installations by month by channel.

Table 4: Download numbers for scalafmt.

Channel	Version	Installations
IntelliJ	v0.2.5	847
	All	3.273
Maven	v0.2.5	788
	All	2.657
Github	v0.2.5	102
	All	929
Sum	v0.2.5	1.737
	All	6.859

editors to run scalafmt no every file save. Amazon famously showed that sales decreased by 1 percent for every 100ms increase in page load time[26]. We believe similar principles apply to scalafmt, every millisecond counts in the user experience. Still, we believe with the current performance scalafmt delivers enough benefits to be useful. As we'll discuss in the following section, our users seem to agree.

5.2 Adoption

Scalafmt has received quite some attention since its release in early March, three months ago. In this section we present the statistics we believe demonstrate that scalafmt is — despite its young age — already proving itself useful for the Scala community. All data points are as of June 9th, 2016.

5.2.1 Installations

Scalafmt has been installed over 6.500 times. Table 4 shows the installation numbers for each official distribution channel. IntelliJ is the Jetbrains plugin repository[8]. Data is not available for how many users built scalafmt from source. The numbers represent absolute download numbers, not unique users. Observe that v0.2.5 was released 22 days ago, meaning it has been installed 80 times on average per day since its release. Considering the

---

[8] See https://plugins.jetbrains.com/plugin/8236?pr=

v0.2.5 installations numbers, it is fair to estimate that scalafmt currently has at least 1.000 users.

Figure 6 shows the growth in installations by month. Observe that growth has doubled with each month. Github represented proportionally many installs in the first month but only represents a small fraction of installation in May. Maven installations quadrupled in May, taking the lead from the IntelliJ plugin in April.

5.2.2   Other

We present interesting data points from a variety of disparate data sources:

- The scalafmt code repository has received contributions from 8 external contributors. Several of these contributions added non-trivial features to scalafmt, including new configuration flags and extensions to the Router.

- 34 unique users, excluding the author, have opened a total of 138 tickets on the scalafmt issue tracker.

- The scalafmt Gitter[9] instant messaging channel has 47 members. The channel is used to informally discuss bugs, new features and more.

- The user documentation website[10] has been visited 5.422 times with an average visit duration of 98 seconds. That makes 138 hours of combined visit time.

---

[9]See https://gitter.im/olafurpg/scalafmt
[10]See http://scalafmt.org

6 Discussion

6.1 Future work

6.2 Conclusion

References

[1] *Akka*. URL: http://akka.io/ (visited on 05/29/2016).

[2] *Apache Spark™ - Lightning-Fast Cluster Computing*. URL: http://spark.apache.org/ (visited on 05/29/2016).

[3] *Bill Gosper*. URL: http://gosper.org/bill.html (visited on 05/31/2016).

[4] Eugene Burmako. "Scala macros: let our powers combine!: on how rich syntax and static types work with metaprogramming". In: *Proceedings of the 4th Workshop on Scala*. ACM. 2013, p. 3.

[5] Eugene Burmako. *scala.meta*. http://scalameta.org/. (Accessed on 06/03/2016). Apr. 2016.

[6] *CodeReviewComments for golang*. https://github.com/golang/go/wiki/CodeReviewComments. (Accessed on 06/01/2016). Oct. 2015.

[7] Daniel Jasper. *clang-format - Automatic formatting for C++*. https://www.youtube.com/watch?v=s7JmdCfI__c. (Accessed on 06/02/2016).

[8] David Copeland Daniel Spiewak. *Scala Style Guide*. http://docs.scala-lang.org/style/. (Accessed on 06/06/2016).

[9] Edsger W. Dijkstra. "A note on two problems in connexion with graphs". In: *Numerische mathematik* 1.1 (1959), pp. 269–271. URL: http://www.springerlink.com/index/uu8608u0u27k7256.pdf (visited on 06/01/2016).

[10] Sébastien Doeraene. *Scala.js Coding Style*. 2015. URL: https://github.com/scala-js/scala-js/blob/master/CODINGSTYLE.md (visited on 05/28/2016).

[11] Stuart Dreyfus. "Richard Bellman on the birth of dynamic programming". In: *Operations Research* 50.1 (2002), pp. 48–51.

[12] *ENSIME*. URL: http://ensime.github.io/ (visited on 05/29/2016).

[13] Olafur Pall Geirsson. *Scalafmt - code formatter for Scala*. Mar. 2016. URL: http://scalafmt.org (visited on 05/29/2016).

[14] Ira Goldstein. *Pretty-printing Converting List to Linear Structure.*
Massachusetts Institute of Technology. Artificial Intelligence
Laboratory, 1973. URL:
http://www.softwarepreservation.net/projects/LISP/MIT/AIM-
279-Goldstein-Pretty_Printing.pdf (visited on 05/29/2016).

[15] *Google C++ Style Guide.* URL:
https://google.github.io/styleguide/cppguide.html (visited on
05/28/2016).

[16] Robert Griesemer. *gofmt - The Go Code Formatter.*
https://golang.org/cmd/gofmt/. (Accessed on 06/01/2016). June
2009.

[17] Sam Halliday. *I don't have time to talk about formatting in code
reviews. I want the machine to do it so I can focus on the design.*
microblog. May 2016. URL:
https://twitter.com/fommil/status/727879141673078785
(visited on 05/29/2016).

[18] Li Haoyi. *sourcecode - Scala library providing "source" metadata to
your program.* https://github.com/lihaoyi/sourcecode. (Accessed
on 06/04/2016). Feb. 2016.

[19] R. W. Harris. "Keyboard standardization". In: 10.1 (1956), p. 37.
URL: http://massis.lcs.mit.edu/archives/technical/western-
union-tech-review/10-1/p040.htm (visited on 05/29/2016).

[20] John Hughes. "The design of a pretty-printing library". In: *Advanced
Functional Programming.* Springer, 1995, pp. 53–96. URL:
http://link.springer.com/chapter/10.1007/3-540-59451-5_3
(visited on 01/06/2016).

[21] Daniel Jasper. *clang-format.* Mar. 2014. URL:
http://llvm.org/devmtg/2013-04/jasper-slides.pdf (visited on
04/20/2016).

[22] Daniel Jasper. *ClangFormat.* 2013. URL:
http://clang.llvm.org/docs/ClangFormat.html (visited on
06/01/2016).

[23] Viktor Klang. *Code style should not be enforced by review, but by
automate rewriting. Evolve the style using PRs against the rewriting
config.* microblog. Feb. 2016. URL:
https://twitter.com/viktorklang/status/696377925260677120
(visited on 05/29/2016).

[24] Jon Kleinberg and Éva Tardos. *Algorithm design.* Pearson Education India, 2006.

[25] Donald E. Knuth and Michael F. Plass. "Breaking paragraphs into lines". In: *Software: Practice and Experience* 11.11 (1981), pp. 1119–1184. URL: http://onlinelibrary.wiley.com/doi/10.1002/spe.4380111102/abstract (visited on 05/31/2016).

[26] Ron Kohavi and Roger Longbotham. "Online experiments: Lessons learned". In: *Computer* 40.9 (2007), pp. 103–105.

[27] Konrad Malawski. *ktoso/sbt-jmh: "Trust no one, bench everything." - sbt plugin for JMH (Java Microbenchmark Harness).* https://github.com/ktoso/sbt-jmh. (Accessed on 06/07/2016).

[28] John McCarthy. "Recursive functions of symbolic expressions and their computation by machine, Part I". In: *Communications of the ACM* 3.4 (1960), pp. 184–195. URL: http://dl.acm.org/citation.cfm?id=367199 (visited on 05/31/2016).

[29] Bob Nystrom. *dart_style - An opinionated formatter/linter for Dart code.* Sept. 2014. URL: https://github.com/dart-lang/dart_style (visited on 06/01/2016).

[30] Bob Nystrom. *The Hardest Program I've Ever Written.* Sept. 2015. URL: http://journal.stuffwithstuff.com/2015/09/08/the-hardest-program-ive-ever-written/ (visited on 04/14/2016).

[31] Martin Odersky and Heather Miller. *The Scala Center.* Mar. 2016. URL: http://www.scala-lang.org/blog/2016/03/14/announcing-the-scala-center.html (visited on 05/29/2016).

[32] Martin Odersky et al. *The Scala language specification.* 2004. URL: http://www-dev.scala-lang.org/old/sites/default/files/linuxsoft_archives/docu/files/ScalaReference.pdf (visited on 05/31/2015).

[33] Bruno CdS Oliveira, Adriaan Moors, and Martin Odersky. "Type classes as objects and implicits". In: *ACM Sigplan Notices.* Vol. 45. 10. ACM. 2010, pp. 341–360.

[34] *OpenJDK: jmh.* http://openjdk.java.net/projects/code-tools/jmh/. (Accessed on 06/07/2016).

[35] Dereck C. Oppen. "Prettyprinting". In: *ACM Trans. Program. Lang. Syst.* 2.4 (Oct. 1980), pp. 465–483. ISSN: 0164-0925. DOI: 10.1145/357114.357115. URL: http://doi.acm.org/10.1145/357114.357115 (visited on 04/18/2016).

[36] Judea Pearl. "Heuristics: intelligent search strategies for computer problem solving". In: (1984). URL: http://www.osti.gov/scitech/biblio/5127296 (visited on 06/01/2016).

[37] Tiark Rompf and Nada Amin. "From F to DOT: Type Soundness Proofs with Definitional Interpreters". In: *arXiv:1510.05216 [cs]* (Oct. 2015). arXiv: 1510.05216. URL: http://arxiv.org/abs/1510.05216 (visited on 05/28/2016).

[38] Matt Russell. *Scalariform*. 2010. URL: http://scala-ide.org/scalariform/ (visited on 05/28/2016).

[39] *sbt - The interactive build tool*. URL: http://www.scala-sbt.org/ (visited on 05/28/2016).

[40] *scala-native/scala-native*. URL: https://github.com/scala-native/scala-native (visited on 05/29/2016).

[41] *Scala.js*. URL: http://www.scala-js.org/ (visited on 05/29/2016).

[42] R. S. Scowen et al. "SOAP—A program which documents and edits ALGOL 60 programs". In: *The Computer Journal* 14.2 (1971), pp. 133–135. URL: http://comjnl.oxfordjournals.org/content/14/2/133.short (visited on 05/29/2016).

[43] *Stack Overflow Developer Survey 2015*. URL: http://stackoverflow.com/research/developer-survey-2015 (visited on 05/29/2016).

[44] S. Doaitse Swierstra and Olaf Chitil. "Linear, bounded, functional pretty-printing". In: *Journal of Functional Programming* 19.01 (Jan. 2009), pp. 1–16. ISSN: 1469-7653. DOI: 10.1017/S0956796808006990. URL: http://journals.cambridge.org/article_S0956796808006990 (visited on 04/20/2016).

[45] Mark Van Den Brand and Eelco Visser. "Generation of formatters for context-free languages". In: *ACM Transactions on Software Engineering and Methodology (TOSEM)* 5.1 (1996), pp. 1–41. URL: http://dl.acm.org/citation.cfm?id=226156 (visited on 01/06/2016).

[46] Philip Wadler. "A prettier printer". In: *The Fun of Programming, Cornerstones of Computing* (2003), pp. 223–243. URL: http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.19.635&rep=rep1&type=pdf (visited on 04/20/2016).

[47] Hyrum Wright et al. "Large-Scale Automated Refactoring Using ClangMR". In: (2013). URL: https://research.google.com/pubs/pub41342.html (visited on 04/21/2016).

[48] Reynold Xin. *Spark Scala Style Guide*. Mar. 2015. URL: https://github.com/databricks/scala-style-guide (visited on 06/01/2016).

[49] Phillip M. Yelland. *A New Approach to Optimal Code Formatting*. Tech. rep. Google, inc., 2016. URL: http://research.google.com/pubs/archive/44667.pdf (visited on 04/20/2016).