



scalafmt: opinionated code formatter for Scala

Ólafur Páll Geirsson

School of Computer and Communication Sciences

Master's Thesis

June 2016

Responsible

Prof. Martin Odersky
EPFL / LAMP

Supervisor

Eugene Burmako
EPFL / LAMP

Abstract

Code formatters bring many benefits to software development such as enforcing a consistent coding style across teams, more effective code reviews and enabling automated large-scale refactoring. Still, code formatters can be tricky to get right. This thesis addresses the problem of developing a code formatter for the Scala programming language. We present `scalafmt`, an opinionated Scala code formatter that captures many popular Scala idioms and coding styles. This thesis introduces the algorithms and tooling we used to implement advanced features such as line wrapping and configurable vertical alignment. We have validated that these techniques work well in practice. `Scalafmt` has been installed over 6,500 times in only 3 months and several popular open-source libraries have chosen to reformat their codebases with `scalafmt`.

Contents

1	Introduction	7
1.1	Contributions	8
2	Background	9
2.1	Scala the programming language	9
2.1.1	Higher order functions	10
2.1.2	Immutability	10
2.1.3	SBT build configuration	11
2.2	scala.meta	12
2.3	Code formatters	14
2.3.1	Natural language	14
2.3.2	ALGOL 60	14
2.3.3	LISP	15
2.3.4	Language agnostic	15
2.3.5	gofmt	17
2.3.6	Scalariform	17
2.3.7	clang-format	18
2.3.8	dartfmt	20
2.3.9	rfmt	21
3	Algorithms	23
3.1	Design	23
3.2	Data structures	24
3.2.1	FormatToken	24
3.2.2	Decision	24
3.2.3	Policy	24
3.2.4	Indent	25
3.2.5	Split	26
3.2.6	State	26
3.3	LineWrapper	27
3.3.1	Router	27
3.3.2	Best-first search	29
3.4	Optimizations	31
3.4.1	dequeueOnNewStatements	31
3.4.2	recurseOnBlocks	33
3.4.3	OptimalToken	34
3.4.4	pruneSlowStates	36
3.4.5	escapeInPathologicalCases	37

3.5	FormatWriter	38
3.5.1	Docstring formatting	39
3.5.2	stripMargin alignment	39
3.5.3	Vertical alignment	40
4	Tooling	43
4.1	Heatmaps	43
4.2	Property based testing	44
4.2.1	Can-format	45
4.2.2	AST integrity	45
4.2.3	Idempotency	46
5	Evaluation	47
5.1	Performance benchmarks	47
5.1.1	Setup	47
5.1.2	Macro benchmark	47
5.1.3	Micro benchmark	49
5.2	Adoption	50
5.2.1	Installations	50
5.2.2	Other	51
6	Discussion	53
6.1	Future work	53
6.2	Conclusion	53

Listings

1	Unformatted code	7
2	Formatted code	7
3	Higher order functions	10
4	Higher order functions expanded	10
5	Manipulating immutable list	11
6	Manipulating mutable list	11
7	SBT project definition	11
8	Parsing different Scala dialects with scala.meta	12
9	Serializing scala.meta trees	13
10	A LISP program	15
11	Gofmt example input/output	17
12	Bin-packing	18
13	No bin-packing	18

14	Unformatted C++ code	19
15	ClangFormat formatted C++ code	20
16	Avoid dead ends	20
17	Line block	21
18	Stack block	21
21	Formatting layout for argument lists	22
19	Line block	22
20	Stack block	22
22	FormatToken definition	24
23	Decision definition	24
24	Policy definition	25
25	Indent definition	25
26	Split definition	26
27	State definition	27
28	Pattern matching on FormatToken	28
29	Unreachable code	29
30	Extracting line number from call site	29
31	Exponential running time	30
32	Two independent statements	31
33	Overeager dequeueOnNewStatements	32
34	recurseOnBlocks example	33
35	OptimalToken definition	34
36	OptimalToken example	35
37	Slow states	36
38	stripMargin example	39
39	Vertical alignment example	40

List of Algorithms

1	Scalafmt best-first search, first approach	30
2	dequeueOnNewStatements optimization	32
3	recurseOnBlocks optimization	34
4	OptimalToken optimization	35
5	pruneSlowStates optimization	36
6	best-effort fallback strategy	38
7	Vertical alignment, simplified algorithm	41
8	AST integrity property	45

List of Figures

1	ClangFormat architecture	19
2	Scalafmt architecture	23
3	Example graph produced by Router	28
4	Example heatmap with 5.121 visisted states	43
5	Example diff heatmap	44
6	Installations by month by channel.	51

List of Tables

1	Results from macro benchmark.	48
2	Percentiles of lines of code per file in micro benchmark.	49
3	Results from micro benchmark.	50
4	Download numbers for scalafmt.	51

1 Introduction

Without code formatters, software developers are responsible for manipulating all syntactic trivia in their programs. What is syntactic trivia? Consider the Scala code snippets in listings 1 and 2.

Listing 1: Unformatted code	Listing 2: Formatted code
<pre>1 object ScalafmtExample { 2 function(arg1, arg2(arg3(3 "String literal"), 4 arg4 + arg5)) 5 } 6 7</pre>	<pre>1 object ScalafmtExample { 2 function(3 arg1, 4 arg2(arg3("String literal"), 5 arg4 + arg5)) 6 } 7</pre>

Both snippets represent the same program. The only difference lies in where the programmer has chosen to break lines. Characters such as spaces and line breaks that do not affect the execution of the program are syntactic trivia. Although syntactic trivia has no meaning for the execution of the program, listing 2 is arguably easier to read, understand and maintain for the software developer. The promise of code formatters is to modify syntactic trivia such that a program that contains style issues, such as in listing 1, turns into a readable and maintainable program, such as in listing 2. Code formatting brings several benefits to software development.

Code formatting enables automated large-scale refactoring. Google used ClangFormat[21], a code formatter, to migrate legacy C++ code to the modern C++11 standard[46]. ClangFormat was used to ensure that the refactored code adhered to Google’s strict C++ coding style[14]. Similar migrations can be expected in the near future for the Scala community once new dialects, such as Dotty[36], gain popularity.

Code formatting is valuable in collaborative coding environments. The Scala.js project[40] has over 40 contributors and the Scala.js coding style[10] — which each Scala.js contributor is expected to know by heart — is defined at a whopping 2,600 word count. Each contributed patch is manually verified against the coding style by the project maintainers. This adds a burden on both contributors and maintainers. Several prominent Scala community members have raised this issue. ENSIME[12] is a popular Scala interaction mode for text editors such as Vim and Emacs. Sam Halliday, an ENSIME maintainer, says “I don’t have time to talk about

formatting in code reviews. I want the machine to do it so I can focus on the design.”[16]. Akka[1] is a popular concurrent and distributed programming library for Scala with over 300 contributors. Viktor Klang, a maintainer of Akka, suggests a better alternative “Code style should not be enforced by review, but by automate rewriting. Evolve the style using PRs against the rewriting config.”[22]. With scalafmt, we hope to relieve Scala developers from the burden of manipulating syntactic trivia so they can instead direct their full attention to writing correct, maintainable and fast code.

1.1 Contributions

The main contributions presented in this thesis are the following:

- language-agnostic algorithms and data structures to implement line wrapping with a maximum line length setting — in our experience, the most challenging part of developing scalafmt — as well as configurable vertical alignment. This work is presented in section 3.
- methods to optimize and test the developed algorithms. This work is presented in section 4.
- practical validation of developed algorithms and methods with scalafmt. The empirical results of this validation are presented in section 5.

2 Background

This chapter explains the necessary background to understand Scala and code formatting. More specifically, we motivate why Scala presents an interesting challenge for code formatters. We go into details on Scala’s rich syntax and popular idioms that introduced unique challenges to the design of scalafmt. We follow up with a history on code formatters that have been developed over the last 70 years. We will see that although code formatters have a long history, a new tradition of optimization based formatters – which scalafmt proudly joins – started only recently in 2013.

2.1 Scala the programming language

Scala[\[31\]](#) is a general purpose programming language that was first released in 2004. Scala combines features from object-oriented and functional programming paradigms, allowing maximum code reuse and extensibility.

Scala can run on multiple platforms. Most commonly, Scala programs compile to bytecode and run on the JVM. With the releases of Scala.js[\[10\]](#), JavaScript has recently become a popular target platform for Scala developers. Even more recently, the announcement of Scala Native[\[39\]](#) shows that LLVM and may become yet another viable target platform for Scala developers.

Scala is a popular programming language. The Scala Center estimates that more than half a million developers are using Scala[\[30\]](#). Large organizations such as Goldman Sachs, Twitter, IBM and Verizon rely on Scala code to run business critical applications. The 2015 Stack Overflow Developer Survey shows that Scala is the 6th most loved technology and 4th best paying technology to work with[\[42\]](#). The popularity of Apache Spark[\[2\]](#), a cluster computing framework for large-scale data processing, has made Scala a language of choice for many developers and scientists working in big data and machine learning.

Scala is a programming language with rich syntax and many idioms. The following chapters discuss in detail several prominent syntactic features and idioms of Scala. Most importantly, we highlight coding patterns that encourage developers to write a single large statement over of multiple small statements. In section [3](#), we will discover that large statements

Listing 3: Higher order functions

```
1 def twice(f: Int => Int) = (x: Int) => f(f(x))
2 twice(_ + 2)(6) // 10
```

Listing 4: Higher order functions expanded

```
1 def twice(f: Function[Int, Int]) =
2   new Function[Int, Int]() { def apply(x: Int) = f.apply(f.apply(x)) }
3 twice(new Function[Int, Int]() { def apply(x: Int) = x + 2 }).apply(6) // 10
```

introduce a challenge to code formatting.

2.1.1 Higher order functions

Higher order functions (HOFs) are a common concept in functional programming languages as well as mathematics. HOFs are functions that can take other functions as arguments as well as return functions as values. Languages that provide a convenient syntax to manipulate HOFs are said to make functions first-class citizens.

Functions are first-class citizens in Scala. Consider listing 3. The method `twice` takes an argument `f`, which is a function from an integer to an integer. The method returns a new function that will apply `f` twice to an integer argument. This small example takes advantage of several syntactic conveniences provided by Scala. For example, in line 2 the argument `_ + 3` creates a new `Function[Int, Int]` object. The function call `f(x)` is in fact sugar for the method call `f.apply(x)` on a `Function[Int, Int]` instance. Listing 4 shows an equivalent program to listing 3, without using syntactic conveniences. Observe that the body of `twice` was expressed as a single statement in line 1 of listing 3 but as two independent statements in listing 4.

2.1.2 Immutability

Functional programming encourages stateless functions which operate on immutable data structures and objects. An immutable object is an object that once initialized, cannot be modified. Immutability offers several

Listing 5: Manipulating immutable list

```
1 val input = List(1, 2, 3)
2 val output = input.map(_ + 1) // List(2, 3, 4)
3                   .filter(_ > 2) // List(3, 4)
```

Listing 6: Manipulating mutable list

```
1 val input = List(1, 2, 3)
2 val output = mutable.ListBuffer.empty[Int] // mutable list
3 input.foreach { elem =>
4   if (elem + 1 > 2) { // filter
5     output += elem + 1 // map
6   }
7 }
8 output // ListBuffer(3, 4)
```

benefits to software development in areas including concurrency and testing. Listing 5 shows an example of manipulating an immutable list. Note that each `map` and `filter` operation creates a new copy of the list with the modified contents. The original list remains unchanged. Listing 6 show the equivalent operation using a mutable list. Observe that listing 5 is a single statement while listing 6 contains multiple statements.

2.1.3 SBT build configuration

SBT[38] is an interactive build tool used by many Scala projects. SBT configuration files are written in `*.sbt` or `*.scala` files using Scala syntax and semantics. Although SBT configuration files use plain Scala, they typically use coding patterns which are different from traditional Scala programs. Listing 7 is an example project definition in SBT. Observe that

Listing 7: SBT project definition

```
1 lazy val core = project
2   .settings(allSettings)
3   .settings(
4     moduleName := "scalafmt-core",
5     libraryDependencies ++= Seq(
6       "com.lihaoyi" %% "sourcecode" % "0.1.1",
7       "org.scalameta" %% "scalameta" % Deps.scalameta))
```

Listing 8: Parsing different Scala dialects with scala.meta

```
1 import scala.meta._
2 dialects.Sbt0137(
3   """lazy val root = project.dependsOn(core)
4     lazy val core = project""").parse[Source] // OK
5 dialects.Sbt0136(
6   """lazy val root = project.dependsOn(core)
7     lazy val core = project""").parse[Source] // Missing blank line
8 // Default dialect, regular Scala compilation unit
9 """lazy val root = project""").parse[Source] // No top-level statements
```

the project is defined as a single statement and makes extensive use of symbolic infix operators. Due to the nature of build configurations, argument lists to can becomes unwieldy long and a single project statement can span over dozens or even hundreds of lines of code.

2.2 scala.meta

Scala.meta^[5] is a metaprogramming toolkit for Scala. Before scala.meta, the state-of-the-art metaprogramming facilities relied on the Scala compiler internals. This had several severe limitations such as too-eager desugaring resulting in loss of syntactic details from the original source code.

Scala.meta was designed to overcome these limitations and offer a more robust platform to develop metaprogramming tools for Scala. Several key features of scala.meta have made it an invaluable companion in the development of scalafmt. Most notably among these features are dialect agnostic syntax trees, syntax tree serialization, high-fidelity parsing and algebraically typed tokens.

Scala.meta provides facilities to tokenize and parse a variety of different Scala dialects. One such dialect is SBT configuration files, discussed in section 2.1.3. SBT adds custom support for top-level statements in *.sbt files. An attempt to parse an SBT file with the Scala compiler parser will result in a parse error. Working around this issue requires either depending on the SBT parser or reimplementing its parsing logic. To add insult to injury, top-level statements must be separated by a blank line if you use an SBT version lower than 0.13.6; a restriction that was lifted in SBT 0.13.7. Listing 8 show how to scala.meta makes it trivial to accommodate this zoo of nuances. The result after parsing is a dialect agnostic scala.meta tree structure.

Listing 9: Serializing scala.meta trees

```
1 import scala.meta._
2 > """ object Main extends App { self =>
3     println(s"Hello $self!") // This is a comment
4 }""".parse[Source].get.structure
5 res0: String = """
6 Source(Seq(Defn.Object(Nil, Term.Name("Main"), Template(Nil, Seq(Ctor.Ref.Name("App
7     ")), Term.Param(Nil, Term.Name("self"), None, None), Some(Seq(Term.Apply(Term.
8     Name("println"), Seq(Term.Interpolate(Term.Name("s"), Seq(Lit("Hello "), Lit
9     ("!")), Seq(Term.Name("self"))))))))))))
10 """)
```

The structure of `scala.meta` trees can be serialized to strip off all syntactic details. Listing 9 shows how to serialize the tree structure of a simple hello world application. For example, observe that the comment has been stripped away. As we discuss in section 4, this feature was instrumental in testing `scalafmt`.

Tree node types in `scala.meta` preserve absolute fidelity with the original source file. This means we can obtain all syntactic details from a tree node such as whether a for comprehension uses parentheses or curly braces as delimiters, whitespace positions, comments and other syntactic trivia. The Scala compiler is infamous for desugaring for-comprehensions into `map/withFilter/flatMap` applications during the parse phase. This made it impossible to implement metaprogramming tasks such as code formatting. High-fidelity parsing in `scala.meta` has been essential for `scalafmt` because we can't lose critical syntactic details such as whether for-comprehensions are used over `flatMap`s.

Tokens in `scala.meta` are strongly typed. Traditional object-oriented libraries treat tokens as a single type with multiple methods such as `isComma/isFor` which returns true if a token instance is a comma or a `for` keyword. However, `scala.meta` leverages algebraic data types in Scala to represent each different kind of token as a separate type. This feature plays nicely with exhaustivity checking in the Scala pattern matcher and enabled design pattern for the *Router* explained in section 3.3.1.

2.3 Code formatters

Code formatting and pretty printing¹ has a long tradition. In this chapter, we look at a variety of tools and algorithm that have been developed over the last 70 years.

2.3.1 Natural language

The science of displaying aesthetically pleasing text predates as early as 1956[18]. The first efforts involved inserting carriage returns in natural language text. Until that time, writers had been responsible for manually providing carriage returns in their documents before sending them off for printing. The motivation was to “save operating labor and reduce human error”. Once type-setting became more commonplace, the methods for breaking lines of text got more sophisticated.

Knuth and Plass developed in 1981 a famous line breaking algorithm[24] for L^AT_EX, a popular typesetting program among scientific academic circles. L^AT_EX is the program that was used to generate this very document. The line breaking problem was the same as in the 60s: how to optimally break a paragraph of text into lines so that the right margin is minimized. The primitive approach is to greedily fit as many words on a line as possible. However, such an approach can produce embarrassingly bad output in the worst case. Knuth’s algorithm uses dynamic programming to find an optimal layout with regards to a fit function that penalizes empty space on the right margin of the paragraph. This algorithm remains a textbook example of an application of dynamic programming[11, 23].

2.3.2 ALGOL 60

Scowen[41] developed SOAP in 1971, a code formatter for ALGOL 60. The main motivation for SOAP was to make it “easier for a programmer to examine and follow a program” as well as to maintain a consistent coding style. This motivation is still relevant in modern software development.

¹ This thesis uses the term *code formatting* over *pretty printing*. According to Hughes[19], pretty printing is a subset of code formatting where the former is only concerned with presenting data structures while the latter is concerned with the harder problem of formatting existing source code — the main topic of this thesis.

Listing 10: A LISP program

```
1 (defun factorial (n)
2   (if (= n 0) 1
3       (* n (factorial (- n 1)))))
```

SOAP did provide a maximum line length limit. However, SOAP would fail execution if the provided line length turned out to be too small. With hardware from 1971, SOAP could format 600 lines of code per minute.

2.3.3 LISP

In 1973, Goldstein[13] explored code formatting algorithms for LISP[27] programs. LISP is a family of programming languages and is famous for its parenthesized prefix notation. Listing 10 shows a program in LISP to calculate factorial numbers. The simple syntax and extensive use of parentheses as delimiters makes LISP programs an excellent ground to study code formatters.

Goldstein presented a *recursive re-predictor* algorithm in his paper. The recursive re-predictor algorithm runs a top-down traversal on the abstract syntax tree of a LISP program. While visiting each node, the algorithm tries to first obtain a *linear-format*, i.e. fit remaining body on a single line, with a fallback to *standard-format*, i.e. each argument is put on a separate line aligned by the first argument. Goldstein observes that this algorithm is practical despite the fact that its running time is exponential in the worst case. Bill Gosper used the re-predictor algorithm to implement GRINDEF[3], one of the first code formatters for LISP.

Goldstein's contributions extend beyond formatting algorithms. Firstly, in his paper he studies how to format comments. Secondly, he presents several different formatting layouts which can be configured by the users. Both relevant concerns for modern code formatters.

2.3.4 Language agnostic

Derek C. Oppen pioneered the work on language agnostic code formatting in 1980[34]. A language agnostic formatting algorithm can be used for a

variety of programming languages instead of being tied to a single language. Users provide a preprocessor to integrate a particular programming language with the algorithm. Oppen’s algorithm runs in $O(n)$ time and uses $O(m)$ memory for an input program of length n and maximum column width m . Besides impressive performance results, Oppen claims that a key feature of the algorithms is its streaming nature. The algorithm prints formatted lines as soon as they are input instead of waiting until the entire input stream has been read. A feature which is typically not a concern for modern code formatters. Moreover, Oppen’s algorithm shares a worrying limitation with SOAP: it cannot handle the case when the line length is insufficiently large.

Mark van der Brand presented a library in 1996 that generates a formatter given a context-free grammar[44]. Beyond the usual motivations for developing code formatters, Brand mentions that formatters “relieve documentation writers from typesetting programs by hand”. The focus on documentation is reflected by the fact that the generated formatter could produce both ASCII formatted code as well as L^AT_EX markup. Since comments are typically not included in a syntax tree, the presented algorithm has an elaborate scheme to infer the location of comments in the produced output. Like Oppen’s algorithm, this library requires the user to plug in a preprocessor to integrate a particular programming language into the Brand’s library. Unlike Oppen’s algorithm, Brand does not consider line length limits in his algorithm.

John Hughes extended on Oppen’s work on language agnostic formatting in term of functional programming techniques[19]. Hughes presented a design of a *pretty-printing* library that leverages combinators with algebraic properties to express formatting layouts. Hughes claims that such a formal approach was invaluable when designing the pretty-printing library, which has seen wise use, including in the Glasgow Haskell compiler. Wadler[45] and Chitil[43] extend on Hughes’s and Oppen’s work in term of performance and programming techniques. However, this branch of work has been limited to printing data structures and not how to format existing source code.

2.3.5 gofmt

`gofmt`[15] is a code formatter for the Go programming language, developed at Google. `gofmt` was released in the early days of Go in 2009 and is noteworthy for its heavy adoption by the Go programming community. Official Go documentation[6] claims that almost all written Go code — at Google and elsewhere — is formatted with `gofmt`. Besides formatting, `gofmt` is used to automatically migrate Go codebases from legacy versions to new source-incompatible releases. However, `gofmt` supports neither a column limit nor an opinionated setting. Line breaks are preserved in the user’s input. For example, listing 11 shows a Go program that uses the same layout as the “unformatted” code (listing 1) in the introduction.

Listing 11: Gofmt example input/output

```
1 package main
2
3 func main() int {
4     function(arg1, arg2(arg3(
5         "String literal"),
6         arg4+arg5))
7 }
```

The output of running `gofmt` through listing 11 is identical to the input. This un-opinionated behavior may be considered desirable by many software developers. However, this thesis is only concerned with opinionated code formatting.

2.3.6 Scalariform

Scalariform[37] was released in 2010 and is a widely used code formatter for Scala. Like `gofmt`, Scalariform does an excellent job of tidying common formatting errors. Moreover, Scalariform supports a variety of configuration options. Scalariform is also impressively fast, it can format large files with over 4.000 lines of code in under 250 milliseconds on a modern laptop. However, Scalariform shares the same limitations with `gofmt`: it lacks a line length and opinionated setting.

Firstly, the line length setting is necessary to implement many popular coding styles in the Scala community. For example, the Spark[47] and Scala.js[10] coding styles have 100 character and 80 character column

limits, respectively. As we see in other code formatters, adding a line length setting is non-trivial and doing so would require a significant redesign of Scalariform.

Secondly, the lack of an opinionated setting makes it impossible to enforce certain coding styles. For example, the Scala.js coding style enforces *bin-packing*, where arguments should be arranged compactly up to the column length limit and indented by 4 spaces. Listings 12 and 13 shows an example of bin packing enabled and disabled, respectively.

Listing 12: Bin-packing	Listing 13: No bin-packing
<pre> 1 // Column 35 2 class Foo(val x: Int, val y: Int, 3 val z: Int) 4 </pre>	<pre> 1 // Column 35 2 class Foo(val x: Int, 3 val y: Int, 4 val z: Int) </pre>

Since Scalariform preserves the line breaking decisions from the input, Scalariform has no setting to convert formatted code like in listing 13 to the code in listing 12.

2.3.7 clang-format

Daniel Jasper triggered a new trend in optimization based coded formatters with the release of *ClangFormat*[20] in 2013. ClangFormat is developed at Google and can format C, C++, Java, JavaScript, Objective-C and Protobuf code. An impressive number of languages. Figure 1 shows the architecture of ClangFormat. The main components are the *structural parser* and the *layouter*.

ClangFormat employs a structural parser to split source code into a sequence of *unwrapped-lines*. An unwrapped line is a statement that should fit on a single line if given sufficient line length. A key feature of unwrapped lines is that they should not influence other unwrapped lines. The parser is lenient and parses even syntactically invalid code. The parsed unwrapped lines are passed onto the layouter.

The ClangFormat layouter uses a novel approach to implement line wrapping. Each line break is assigned a penalty according to several rules such as nesting and token type. At each token, the layouter can choose to continue on the same line or break. This forms an acyclic weighted directed




Figure 1: ClangFormat architecture

Listing 14: Unformatted C++ code

```
1 int main(int argc,char const*argv[]) { Defn.Object( Nil, "ClangFormat", Term.Name("
  State"), Foo.Bar( Template( Nil, Seq( Ctor.Ref.Name("ClangLogger")), Term.Param(
    Nil, Name.Anonymous(), None, None)) ), Term.Name("clang-format") ); }
```

graph with tokens representing vertices and splits (e.g., space, no space or line break) representing edges. The first token of an unwrapped line is the root of the graph and all paths end at the last token of the unwrapped line. The layouter uses Dijkstra’s[9] shortest path algorithm to find the layout that has the lowest penalty. To obtain good performance, the layouter uses several domain specific optimizations to minimize the search space.

Despite supporting an impressive number of programming languages, ClangFormat does not leverage the language agnostic formatting techniques described section 2.3.4. Support for each language has been added as ad-hoc extensions to the ClangFormat parser and layouter. ClangFormat supports a variety of configuration options, including 6 out-of-the-box styles based on coding styles from Google, LLVM and other well-known organizations.

A notable feature of ClangFormat is that it’s opinionated. ClangFormat produces well-formatted output for even the most egregiously formatted input. Listing 14 shows an offensively formatted C++ code snippet. Listing 15 shows the same snippet after being formatted with ClangFormat. ClangFormat is opinionated in the sense that it does not respect the user’s line breaking decisions. This feature makes it possible to

Listing 15: ClangFormat formatted C++ code

```
1 int main(int argc, char const *argv[]) {  
2     Defn.Object(nil, "ClangFormat", Term.Name("State"),  
3         Foo.Bar(Template(nil, Seq(Ctor.Ref.Name("ClangLogger")),  
4             Term.Param(nil, Name.Anonymous(), None, None))),  
5         Term.Name("clang-format"));  
6 }
```

Listing 16: Avoid dead ends

```
1 // Column 35 |  
2 function(  
3     firstCall(a, b, c, d, e),  
4     secondCall("long argument string"));
```

ensure that all code follows the same style guide, regardless of author.

2.3.8 dartfmt

Dartfmt[28] was released in 2014 and follows the optimization based trend initiated by ClangFormat. Dartfmt is a code formatter for the Dart programming language, developed at Google. Like ClangFormat, dartfmt has a line length setting and is opinionated. Bob Nystrom, the author of dartfmt, discusses the design of dartfmt in an excellent post[29] on his blog. In his post, Nystrom argues that the design of a code formatters is significantly complicated by a column limit setting. The line wrapping algorithm in dartfmt employs a *best-first search*[35], a minor variant of the shortest path search in ClangFormat. As with ClangFormat, a range of domain-specific optimizations were required to make the search scale for real-world code. Listing 16 shows an example of such an optimization, *avoiding dead ends*. The snippets exceeds the 35 character column limit. A plain best-first search would perform a lot of redundant search inside the argument list of `firstCall`. However, `firstCall` already fits on a line and there is no need to explore line breaks inside its argument list. The dartfmt optimized search is able to eliminate such dead ends and quickly figure out to break before the "long argument string" literal.

2.3.9 rfmt

The most recent addition to the optimization based formatting trend is `rfmt`[48], a code formatter for the statistical programming environment *R*. The formatter was released in 2016 – after the background work on this thesis started – and like its forerunners is also developed at Google. `rfmt` makes an interesting contribution in that it combines the algebraic combinator approach from Hughes[19] and the optimization based approach from L^AT_EX and ClangFormat.

The algebraic combinator approach makes it easy to express a variety of formatting layouts. `rfmt` uses 6 layout combinators or *blocks* as they are called in the report. The blocks are the following:

- *TextBlock*(*txt*): unbroken string literal.
- *LineBlock*(*b*₁, *b*₂, . . . , *b*_{*n*}): horizontal combination of blocks.
- *StackBlock*(*b*₁, *b*₂, . . . , *b*_{*n*}): vertical combination of blocks.
- *ChoiceBlock*(*b*₁, *b*₂, . . . , *b*_{*n*}): selection of a best block.
- *IndentBlock*(*n*, *l*): indent block *b* by *n* spaces.
- *WrapBlock*(*b*₁, *b*₂, . . . , *b*_{*n*}): Fit as many blocks on each line as possible, break when the column limit is exceeded and align by the first character in *b*₁.

We’ll use an example to show how these relatively few combinators allow an impressive amount of flexibility. Listings 17 and 18 shows two different layouts to format an argument list.

Listing 17: Line block	Listing 18: Stack block
<pre> 1 // Column 35 2 function(argument1, argument2, 3 argument3, argument4, 4 argument5, argument6) 5 </pre>	<pre> 1 // Column 35 2 function(3 argument1, argument2, argument3, 4 argument4, argument5, argument6 5) </pre>

In this case, we prefer the line block from listing 17 since it requires fewer lines. However, our preference changes if the function name is longer as is shown in listings 19 and 20.

Listing 21: Formatting layout for argument lists

```
ChoiceBlock(LineBlock(LineBlock(TextBlock(f), TextBlock("(")),
    WrapBlock(a1, ... , am),
    TextBlock(")"),
    StackBlock(LineBlock(TextBlock(f), TextBlock("(")),
        IndentBlock(4, WrapBlock(a1, ... , am)),
        TextBlock(")")).
```

Listing 19: Line block

```
1 // Column 35
2 functionNameIsLonger(argument1,
3     argument2,
4     argument3,
5     argument4,
6     argument5,
7     argument6)
8
```

Listing 20: Stack block

```
1 // Column 35
2 functionNameIsLonger(
3     argument1, argument2, argument3,
4     argument4, argument5, argument6
5 )
6
7
8
```

Here, we clearly prefer the stack block in listing 20. Listing 21 shows how we use the 6 fundamental blocks in the `rfmt` combinator algebra to express the choice between these two formatting layouts. The variable f denotes the function name and a_1, \dots, a_m denotes the argument list. Observe that listing 21 does not express how to find the optimal layout.

To find an optimal layout, `rfmt` employs a smart dynamic programming trick. First, it is possible to enumerate all layout combinations like the re-predictor algorithm does in section 2.3.3. This leads to exponential growth which turns out to be a problem for some cases. Dynamic programming comes to the rescue by allowing us to reuse partial solutions. Instead of re-calculating the layout cost at each (starting column, block) pair, we store the result in an associative array keyed by the starting column. However, it turns out that this can still be inefficient in terms of memory and speed². To overcome this limitation, Yelland – the `rfmt` author – presents an indexing scheme that makes it possible to extrapolate the layout cost even for missing keys. We refer to the original paper[48] for details. This novel approach enables `rfmt` to format even the most pathologically nested code in near instant time.

² In fact, ClangFormat started with a similar approach, as explained in this[7] video recording, but then switched to Dijkstra’s shortest path algorithms (which in itself is another form of dynamic programming).




Figure 2: Scalafmt architecture

3 Algorithms

This chapter describes how scalafmt formats Scala code. We will see that scalafmt’s design is inspired by ClangFormat and dartfmt. However, we believe our design makes a valuable contribution in that it leverages functional programming principles to maximise code reuse and extensibility (TODO, really??).

3.1 Design

Figure 2 shows a broad architectural overview of scalafmt. First, scalafmt parses a source file using `scala.meta`. Next, we feed a sequence of *FormatToken* data types into a *LineWrapper*. The *LineWrapper* uses a *Router* to construct a weighted directed graph and run a best-first search to find an optimal formatting layout for the whole file. Finally, the *LineWrapper* feeds a sequence of *Split* data types into the *FormatWriter*, which constructs a new reformatted source file. The following sections explain these data types and abstractions in detail.

3.2 Data structures

Scalafmt leverages a few carefully designed data structure to allow an implementation that emphasizes correctness and maintainability.

3.2.1 FormatToken

A *FormatToken* is a pair of two non-whitespace tokens. Listing 22 shows the definition of the `FormatToken` data type.

Listing 22: `FormatToken` definition

```
1 case class FormatToken(left: Token, right: Token, between: Vector[Whitespace])
```

As shown in the architecture overview in figure 2, each token except the beginning and end of file tokens appear twice in the sequence of `FormatTokens`: once as the `left` member and once as the `right` member. In a nutshell, the job of the `LineWrapper` is to convert each `FormatToken` into a *Split*

3.2.2 Decision

A `Decision` is a pair of a `FormatToken` and a sequence of `Splits`. Listing 23 shows the Definition of decision.

Listing 23: `Decision` definition

```
1 case class Decision(formatToken: FormatToken, splits: Seq[Split])
```

The *splits* member represents the possible splits that the `LineWrapper` can choose for *formatToken*.

3.2.3 Policy

A *Policy* is an enforced formatting layout over a region. Listing 24 shows the definition of `Policy`.

Listing 24: Policy definition

```
1 case class Policy(f: PartialFunction[Decision, Decision], expire: Token)
```

A Policy is a partial function that should be applied to future Decisions up until the *expire* token. Policies easily compose using the Scala standard library `orElse` and `andThen` methods on `PartialFunction`³. Policies enable a high-level way to express arbitrary formatting layouts over a region of code.

3.2.4 Indent

An *Indent* describes indentation over a region of code.

Listing 25: Indent definition

```
1 sealed abstract class Length
2 case class Num(n: Int) extends Length
3 case object StateColumn extends Length
4
5 case class Indent[T <: Length](length: T, expire: Token, inclusive: Boolean)
```

Listing 25 shows the definition of *Indent* along with the algebraic data type *Length*. *Length* can either be *Num(n)* where *n* represents an explicit number of spaces to indent by or *StateColumn* which is a placeholder for the number of spaces required to vertically align by the current column. *Indent* is type parameterized by *Length* so that, at some point, we can replace *StateColumn* placeholders with *Nums* to obtain a concrete number. For example, given a `scala.meta` tree `expr`, the definition `Indent(Num(2), expr.tokens.last, inclusive=true)` increases the indentation level by 2 spaces up to and including the last token of `expr`. The `inclusive` member is set to false when the indentation should expire before the expire token, for example in a block wrapped by curly braces, since the closing curly brace should not be indented by 2 spaces. The *StateColumn* placeholder is required to allow memoization of *Splits*, which is critical for performance reason as explained in section 3.3.1 on the *Router*.

³ Fun fact. Careful eyes will observe that *Policy* is in fact a monoid with the empty partial function as identity and function composition as associative operator.

3.2.5 Split

A *Split* represents a (possibly empty) whitespace character to be inserted between two non-whitespace tokens. Listing 26 shows the rather intricate definition of the *Split* data type⁴.

Listing 26: Split definition

```
1 case class Split(modification: Modification,  
2                 cost: Int,  
3                 policy: Policy,  
4                 optimalAt: Option[OptimalToken],  
5                 indents: Vector[Indent[Length]])(  
6   implicit val line: sourcecode.Line)
```

The *Split* data type went through several generations of design before reaching its current structure. Each member serves an important role. The most important member of the *Split* type is the *modification*. A *modification* must be one of *NoSplit*, *Space* and *Newline*. The *cost* member represents the penalty for choosing this split. The *optimalToken* member enables an optimization explained in section 3.4.3. The *line* member allows a powerful debugging technique explained in section 3.3.1. The *policy* and *indents* members are explained in sections 3.2.3 and 3.2.4, respectively.

3.2.6 State

A *State* is a partial formatting solution during the best-first search. Listing 27 shows the definition of the *State* class and companion object.

⁴ For clarity reasons, a few less important members have been removed from the actual *Split* definition.

Listing 27: State definition

```
1 case class State(splits: Vector[Split],
2                 totalCost: Int,
3                 policies: Vector[Policy],
4                 indents: Vector[Indent[Num]],
5                 indentation: Int,
6                 column: Int,
7                 formatOff: Boolean) extends Ordered[State] {
8
9   def compare(that: State): Int
10 }
11
12 object State {
13   def nextState(currentState: State, formatToken: FormatToken, split: Split): State
14 }
```

Observe the similarity of State and Split. A State contains various summaries calculated from the `splits` vector. The summaries are necessary for performance reasons in the best-first search. Observe that the indents are type parameterized by `Num`, meaning they only contain concrete indentations and no `StateColumn` indents. The `indentation` member is the sum of all currently active indents and `column` represents how many characters have been consumed since the last newline. The State class extends the `Ordered` trait to allow for efficient polling from a priority queue. The `compare` method orders States firstly by their `totalCost` member, secondly by `splits.length` – how many FormatTokens have been formatted – and finally breaking ties by the `indentation`. The method `nextState` calculates a penalty for characters that overflow the column limit and prepares The method is implemented as efficiently as possible since the method is on a hot path in the best-first search.

3.3 LineWrapper

The LineWrapper is responsible for turning FormatTokens into Splits. To accomplish this, the LineWrapper employs a *Router* and a *best-first search*.

3.3.1 Router

The Router’s role is to produce a Decision given a FormatToken. Figure 3 shows all possible formatting layout for the small input `val x = y + z`. In this figure, the Router is the planner that has chosen to open up multiple




Figure 3: Example graph produced by Router

branches at = and + and only one branch for the remaining tokens. This is no easy task since a `FormatToken` can be any pair of two tokens. How do we go about implementing a Router?

The Router is implemented as one large pattern match on a `FormatToken`. Listing 28 shows how we can pattern match on a `FormatToken` and produce Splits.

Listing 28: Pattern matching on `FormatToken`

```

1 formatToken match {
2   case FormatToken(_, _: Keyword, _) => Seq(Split(Space, 0))
3   case FormatToken(_, _: '=', _)    => Seq(Split(Space, 0))
4   case FormatToken(_, '(', _)        => Seq(Split(Space, 0)
5                                           Split(Newline, 1))
6   // ...
7 }

```

The pattern `_: '='` matches a `scala.meta` token of type `'='`. The underscore `_` ignores the underlying value. `Keyword` is a super-class of all `scala.meta` keyword token types. Now, a good observer will notice that this pattern match can quickly grow unwieldy long once you account for all of Scala's rich syntax. How does this solution scale? Also, once the match grows bigger how can we know from which case each `Split` origins? It turns out that Scala's pattern matching and `scala.meta`'s algebraically typed tokens are able to help us.

The Scala compiler can statically detect unreachable code. If we add a case that is already covered higher up in the pattern match, the Scala compiler issues a warning. For example, listing 29 shows an example where the compiler issues a warning.

Listing 29: Unreachable code

```
1 formatToken match {  
2   case FormatToken(_, _: Keyword) => Seq(Split(Space, 0))  
3   // ...  
4   case FormatToken(_, _: 'else') => Seq(Newline(), 0) // Unreachable code!  
5 }
```

Here, we accidentally match on a `FormatToken` with an `else` keyword on the right which will never match because we have a broader match on a `Keyword` higher up. In this small example, the bug may seem obvious but once the Router grows bigger the compiler becomes unmissable. However, this still leaves us with the second question of finding the origin of each `Split`. Scala macros^[4] and implicits^[32] come to the rescue.

The source file line number of where a `Split` is instantiated is automatically attached on each `Split`. Remember in listing 26 that the `Split` case class had an implicit member of type `sourcecode.Line`. `Sourcecode`^[17] is a tiny Scala library to extract source code metadata from your programs. The library leverages Scala macros and implicits to unobtrusively surface useful information such as line number of call sites. Listing 30 shows how this works.

Listing 30: Extracting line number from call site

```
1 Split(Space, 0) /* expands into */ Split(Space, 0)(sourcecode.Line(1))
```

When a `sourcecode.Line` not passed explicitly as an argument to `Split`'s constructor, the Scala compiler will trigger its implicit search to fill the missing argument. The `sourcecode.Line` companion contains an implicit macro that generates a `Line` instance from an extracted line number. Take a moment to appreciate how these two advanced features of the Scala programming language enable a very powerful debugging technique. The `scalafmt` router implementation contains 88 cases and spans over 1.000 lines of code. The ability to trace the origin of each `Split` has been indispensable in the development of the Router.

3.3.2 Best-first search

The Decisions from the Router produce a directed weighted graph, as demonstrated in figure 3. To find the optimal formatting layout, our challenge is to find the cheapest path from the root node to a final node.

The best-first[35] algorithm is an excellent fit for the task.

Best-first search is a graph search algorithm to efficiently traverse a directed weighted graph. The objective is reach the final token and once we reach there, we terminate the search because we’re guaranteed no other solution is better. Algorithm 1 shows a first attempt⁵ to adapt a best-first search algorithm to the data structures and terminologies introduced so far. In the best case, the search always chooses the cheapest splits and the

Algorithm 1: Scalafmt best-first search, first approach

```

1  /** @returns Splits that produce and optimal formatting layout */
2  def bestFirstSearch(formatTokens: List[FormatTokens]): List[Split] = {
3    val Q = mutable.PriorityQueue(State.init(formatTokens.head))
4    while (Q.nonEmpty) {
5      val currentState = Q.pop
6      if (currentState.formatToken == formatTokens.last) {
7        return currentState.splits // reached the final state.
8      } else {
9        val splits = Router.getSplits(currentState.formatToken)
10       splits.foreach { split =>
11         Q += State.nextState(currentState, split)
12       }
13     }
14   }
15   // Error: No formatting solution found.
16   ???
17 }

```

algorithm runs in linear time. Observe that the router is responsible for providing well-behaved splits so that we never hit on the error condition after the while loop. Excellent, does that mean the search is complete? Absolutely not, this implementation contains several serious performance issues.

Algorithm 1 is exponential in the worst case. For example, listing 31 shows a tiny input that triggers the search to explore over 8 million states.

Listing 31: Exponential running time

```

1  // Column 60 |
2  a + b + c + d + e + f + g + h + i + j + k + l +
3  m + n + o + p + q + r + s + t + v + w + y +
4  // This comment exceeds column limit, no matter what path is chosen.
5  z

```

⁵ Unfortunately, we make heavy use of mutation since graph search algorithms typically don’t lend themselves well to functional programming principles.

Even if we could visit 1 state per microsecond — reality is closer to 10 states/microsecond — the search will take almost 1 second to complete. This is unacceptable performance to format only 2 lines of code. Of course, we could special-case long comments, but that would only provide us a temporary solution. Instead, like with ClangFormat and dartfmt, we must apply several domain specific optimizations. In the following section, we discuss the optimizations that have shown to work well for scalafmt.

3.4 Optimizations

This section explains the most important domain-specific optimizations that were required to get acceptable performance for scalafmt. We will see that some optimizations are quite ad-hoc and require some creative workarounds.

3.4.1 dequeueOnNewStatements

Once the search reaches the beginning of a new statement, empty the priority queue. Observe that the formatting layout for each statement is independent from the formatting layout of the previous statement. Consider listing 32.

Listing 32: Two independent statements

```
1 // Column 60 |
2 statement1(argument1, argument2, argument3, argument5, argument6)
3 statement2(argument1, argument2, argument3, argument5, argument6)
```

Both statements exceed the column limit, which means that the search must back-track to some extent. However, once the search reaches `statement2` we have already found an optimal formatting layout for `statement1`. When we start backtracking in `statement2`, there is no need to explore alternative formatting layouts for `statement1`. Instead, we can safely empty the search queue once we reach the `statement2` token.

The `dequeueOnNewStatements` optimization is implemented by extending algorithm 1 with an if statement. Algorithm 2 shows a rough sketch of how this is done. With an empty queue, we ensure the search backtracks only as far back as is needed. The `statementStarts` variable contains all tokens

Algorithm 2: dequeueOnNewStatements optimization

```
1 // ...
2 val statementStarts: Set[Token]
3 while (Q.nonEmpty) {
4     val currentState = Q.pop
5     if (statementStarts.contains(currentState.formatToken.left)) {
6         Q.dequeueAll // empty search queue
7     }
8     // ...
9 }
```

that begin a new statement. To collect those tokens, we traverse the syntax tree of the input source file and select the first tokens of each statement of a block, each case in a partial function, enumerator in a for comprehension and so forth. The actual implementation is quite elaborate and is left out of this thesis for clarity reasons. Unfortunately, our optimization has one small problem.

Algorithm 2 may dequeue too eagerly inside nested scopes, leading the search to hit the error condition. Listing 33 shows an example where this happens.

Listing 33: Overeager dequeueOnNewStatements

```
1 // Column 50 |
2 function1(argument1, { case 'argument2' => 11 }, argument3 // forced newline
3                 argument4)
```

Remember that each case of a partial function starts a new statement. The `dequeueOnNewStatements` optimization will dequeue the queue on the first state that reaches the `case` token. In this example, the first state to reach the `case` token will have a strict Policy that disallows newlines up until the closing parenthesis. However, we must insert a newline after the comment. This causes the search to terminate too early and reach the error condition. By inspecting where this problem occurred, we came up with a simple rule to identify regions where the `dequeueOnNewStatements` optimization should be disabled. The simple rule is to never run `dequeueOnNewStatements` inside a pair of parentheses. In section 4, we discuss techniques we used to be confident that this rule indeed works as intended. In the following section (3.4.2) we explain the `recurseOnBlocks` optimization, which allows us to reenable `dequeueOnNewStatements` for selected regions inside parentheses.

3.4.2 recurseOnBlocks

If the `dequeueOnNewStatements` optimization is disabled and we start a new block delimited by curly braces, recursively run the best-first search inside the block. The intuition here is that by recursively running the best-first search, we keep the priority queue small at each layer of recursion. This allows us to run aggressive optimizations such as `dequeueOnNewStatements`.

The `recurseOnBlocks` optimization enables scalafmt to handle idiomatic Scala code where large bodies of higher order functions and blocks are passed around as arguments. Remember from section 2.1 that Scala makes it syntactically convenient to in higher order functions as arguments to other functions. Listing 34 shows an example where this happens and we trigger the `recurseOnBlocks` optimization.

Listing 34: `recurseOnBlocks` example

```
1 function(argument1, { higherOrderFunctionArgument =>
2   statement1
3   // ...
4   statementN
5 })
```

The `dequeueOnNewStatements` optimization is disabled inside argument list. The priority queue grows out bounds because the higher order function can have an arbitrary number of statements.

To implement the `recurseOnBlocks` optimization, we add an extension to algorithm 1. Algorithm 3 shows a rough sketch of how `recurseOnBlocks` is implemented. We change the signature to accept a starting State and token where we stop the search. Observe that we guard against infinite recursion by not making a recursive call on `start.formatToken`. With `recurseOnBlocks` and `dequeueOnNewStatements`, we have solved most problems caused by independent statements affecting the formatting layouts of each other. Next, we leverage recursion again to help the search queue stay small.

Algorithm 3: recurseOnBlocks optimization

```
1 def bestFirstSearch(start: State, stop: Token): List[Split] = {
2   val Q = mutable.PriorityQueue(start)
3   while (Q.nonEmpty) {
4     val currentState = Q.pop
5     if (currentState.formatToken.left == stop) {
6       return currentState
7     } else if (currentState.formatToken != start.formatToken &&
8               currentState.formatToken.left.isInstanceOf['{']) {
9       bestFirstSearch(currentState, closingCurly(currentState.formatToken.left))
10    }
11    // ...
12  }
13 }
```

Listing 35: OptimalToken definition

```
1 case class OptimalToken(token: Token, killOnFail: Boolean = false)
```

3.4.3 OptimalToken

An OptimalToken is a hint from a Split to the best-first search that enables the search to early eliminate competing Splits. Recall from listing 26 that a Split has an `optimalToken` member. Listing 35 shows the definition of OptimalToken. When the best-first search encounters a Split with a defined OptimalToken, the best-first search makes an attempt to reach that token with a budget of 0 cost. If successful, the search can eliminate the competing Splits. If unsuccessful and the `killOnFail` member is true, the best-first search eliminates the Split. Otherwise, the best-first search continues as usual.

By eliminating competing branches, we drastically minimize the search space. Listing 36 shows an example where the OptimalToken can be applied. Scalafmt supports 4 different ways to format call-site function applications. This means that there will be 4^N number of open branches when the search reaches `UserObject N`. To overcome this issue, we define an OptimalToken at the closing parenthesis. The best-first search successfully fits the argument list of each `UserObject` on a single line, and eliminates the 3 other competing branches. This makes the search run in linear time as opposed to exponential.

Listing 36: OptimalToken example

```

1 // Column 50 |
2 Database(
3   UserObject(name1, age1),
4   UserObject(name2, age2),
5   // ...
6   UserObject(nameN, ageN) // comment will always exceed column limit
7 )

```

To implement the `OptimalToken` optimization, we add an extension to algorithm 3. Algorithm 4 sketches how the extension works. The

Algorithm 4: OptimalToken optimization

```

1 def bestFirstSearch(start: State, stop: Token, maxCost: Int): List[Split] = {
2   // ...
3   val splits = Router.getSplits(currentState.formatToken)
4   var optimalFound = false
5   splits.withFilter(_.cost < maxCost).foreach { split =>
6     val nextState = State.nextState(currentState, split)
7     split.optimalToken match {
8       case Some(OptimalToken(expire, killOnFail)) =>
9         val nextNextState = bestFirstSearch(nextState, expire, maxCost = 0)
10        if (nextNextState.expire == expire) {
11          optimalFound = true
12          Q += nextNextState
13        } else if (!killOnFail) {
14          Q += nextState
15        }
16      case _ if !optimalFound =>
17        Q += nextState
18    }
19  }
20  // ...
21 }

```

`bestFirstSearch` method has a new `maxCost` parameter, which is the highest cost that a new splits can have. Next, if a Split has defined an `OptimalToken` we make an attempt to format up to that token. If successful, we update `optimalFound` variable to eliminate other Splits from being added to the queue. If unsuccessful and `killOnFail` is true, we eliminate the Split that defined the `OptimalToken`. A straightforward extension to this algorithm would be to add a `maxCost` member to the `OptimalToken` definition from listing 35. However, this has not been necessary for scalafmt.

3.4.4 pruneSlowStates

The `pruneSlowStates` is a optimization that eliminates states that progress slowly. A state progresses slowly if it visits a token after other equally or less expensive states. The insight is that if two equally expensive states visit the same token, the first state to visits that token typically produces a better formatting layout.

By eliminating slow states, we obtain a better formatting output in addition to minimizing the search space. Listing 37 shows two equally expensive formatting solutions where one solution is fast and other is slow.

Listing 37: Slow states

```
1 // Column 30          |
2
3 // Fast state
4 a + b + c + d + e + f + g +
5 h + i + j
6 // slow state
7 a + b + c +
8 d + e + f + g + h + i + j
```

Of course, the line break after `g +` could be more expensive than the line break after `c +`. Instead, the Router is free to assign an identical cost to both line breaks and let `pruneSlowStates` take care of eliminating the slow state.

The `pruneSlowStates` is implemented as a extension to algorithm 4. Algorithm 5 shows a rough sketch of how the extension works.

Algorithm 5: `pruneSlowStates` optimization

```
1 // ...
2 val fastStates: mutable.Map[FormatToken, State]
3 while (Q.nonEmpty) {
4   val currentState = Q.pop
5   if (fastStates.get(currentState.formatToken)
6       .exists(_.cost <= currentState.state) {
7     // do nothing, eliminate currentState because it's slow.
8   } else {
9     if (!fastStates.contains(currentState.formatToken)) {
10      // currentState is the fastest state to reach this token.
11      fastStates.update(currentState.formatToken, currentState)
12    }
13    // continue with algorithm
14  }
15 }
```

Observe that this algorithm is transparent to the Router. No special annotations are required from Splits.

3.4.5 `escapeInPathologicalCases`

Alas, despite our best efforts to keep the search space small, some inputs can still trigger exponential running times. The `escapeInPathologicalCases` optimization is our last resort to handle such pathological inputs. How do we detect that the search has encountered such a challenging input?

We detect the search space is growing out of bounds by tallying the number of visits per token. If we visit the same token N times, we can estimate the current branching factor to be around $\log_2(N)$. In `scalafmt`, we tune N to be 256 so that the best-first search can split into two or more paths for up to 8 tokens. When a token has been visited more than 256 times, we trigger the `escapeInPathologicalCases` optimization. In the following paragraphs section, we present two alternative fallback strategies: *leave unformatted* and *best-effort*.

The simplest and most obvious fallback strategy is to leave the pathologically nested code unformatted. This can be implemented by backtracing to the first token of the current statement and then reproduce the formatting input up to the last token of said statement. This method is guaranteed to run linearly to the size of the input. The responsibility is left to the software developer to manually format her code, removing all the benefits of code formatting. However, in some cases the software developer may prefer to get a decent yet suboptimal formatting output.

The best-effort fallback strategy applies heuristics to give a decent but suboptimal formatting output. When a token is visited for the 256th time, we select two candidate states from the search queue and eliminate all other states. The first state is the fastest state — the state that has reached furthest into the token stream — that is not bound a prohibitive single line policy. Single line policies are policies that eliminate newline Splits. The second state is the current state — the slow state that visited the token for the 256th time. The intuition is that the fast state has good formatting output so far but is stuck on a challenging token for some reason. The slow may paid a hefty penalty early causing it to move slowly but maybe the early penalty will yield a better output in the end.

Algorithm 6 shows an example of the best-effort strategy can be implemented as an extension to algorithm 1. The `isSafe` method on `State`

Algorithm 6: best-effort fallback strategy

```

1  var fastestState: State
2  val visits: mutable.Map[FormatToken, Int].withDefaultValue(0)
3  while (Q.nonEmpty) {
4    val currentState = Q.pop
5    visits.update(currentState.formatToken, 1 + visits(currentState.formatToken))
6    if (currentState.length > fastestState.length && currentState.isSafe) {
7      fastestState = currentState
8    }
9    if (visits(currentState.formatToken) == MAX_VISITS_PER_TOKEN) {
10     Q.dequeueAll
11     Q += fastestState
12     Q += currentState
13     visits.clear()
14   } else {
15     // continue with algorithm
16   }
17 }

```

returns true if the state contains prohibitive policies, derived from annotated metadata in Splits from the Router. Observe that this algorithm will reapply the best-effort fallback until the search reaches the final token. In `scalafmt`, we bound the number of this can happen with a final fallback to the unformatted strategy.

The unformatted and best-effort fallback strategies offer different trade-offs. The unformatted strategy works well in a scenario where a software developer is available to manually fix formatting errors. The best-effort strategy works well on computer generated code where just a modicum of formatting still greatly aid the legibility of the code. Unfortunately, as we discuss in section 4.2, we struggled to guarantee idempotency using the best-effort strategy. This limitation renders the best-effort strategy useless in environments where code formatters are used to enforce a consistent coding style across a codebase.

3.5 FormatWriter

Recall from figure 2, the `FormatWriter` receives splits from the best-first search and produces the final output presented to the user. In addition to reifying Splits, the `FormatWriter` runs three post-processing steps:

docstring formatting, vertical alignment and stripMargin alignment.

3.5.1 Docstring formatting

Docstrings are used by software developers to document a specific part of code. Like in Java, docstrings in Scala start with the `/**` pragma and end with `*/`. However, unlike in Java, the Scala community is split on whether new lines inside docstrings should align by the first or the second asterisk. The official Scala Style Guide[8] dictates that new lines should align by the second asterisk while the Java tradition is to align by the first asterisk. The Scala.js[10] and Spark[47] style guides follow the Java convention. To accommodate all needs, scalafmt allows the user to choose either style. To enforce that the asterisks are aligned according to the user's preferences, the `FormatWriter` rewrites docstring tokens. This is implemented with simple regular expressions.

3.5.2 stripMargin alignment

The Scala standard library adds a `stripMargin` extension method on strings. The method helps Scala developers write multiline interpolated and regular string literals. Listing 38 shows an example usage of the `stripMargin` method.

Listing 38: `stripMargin` example

```
1 object StripMarginExample {  
2   """Multiline string are delimited by triple quotes in Scala.  
3   |You can write as many lines as you want.""".stripMargin  
4 }
```

After calling the method, the indentation and `|` character on line 3 are conveniently removed. However, the hard-fought indentation on the pipe can easily be lost when the string is moved up or down a scope during refactoring. Scalafmt can automatically fix this issue. In the `FormatWriter`, scalafmt rewrites string literals to automatically align the `|` characters with the opening triple quotes `"""`. This setting is disabled by default since scalafmt requires semantic information to confidently determine if the `stripMargin` invocation calls the standard library method or a user-defined method.

3.5.3 Vertical alignment

It turns out that vertical alignment is incredibly popular in the Scala programming community. When lines are vertically aligned when similar operators from a sequence of lines align together. Listing 39 shows an example of vertical alignment.

Listing 39: Vertical alignment example

```
1 object VerticalAlignment {
2   x match {
3     // Align by => and -> and //
4     case 1 => 1 -> 2 // first
5     case 11 => 11 -> 22 // second
6
7     // Blank lines separate alignment blocks.
8     case ignoreMe => 111 -> 222
9   }
10
11   def name = column[String]("name")
12   def status = column[Int]("status")
13   val x = 1
14   val xx = 22
15
16   libraryDependencies ++= Seq(
17     "org.scala-lang" % "scala-compiler" % scalaVersion.value,
18     "com.lihaoyi" %% "sourcecode" % "0.1.1"
19   )
20 }
```

The lack of vertical alignment in the initial release of scalafmt was a hindrance for user-adoption. Configurable vertical alignment was added to the 0.2 release.

Vertical alignment is implemented in the FormatWriter as an extension to the reification of Splits. Instead of reifying a Space into a single space literal, the FormatWriter builds a `tokenAligns: Map[Split, Int]` which specifies the number of additional spaces to add for each Space. Listing 7 shows a simplification of how the actual algorithm⁶ constructs the `tokenAligns` map. In a nutshell, the algorithm builds blocks of lines that can be vertical aligned. Blocks are separated by blank lines or mismatching candidates in subsequent lines. As demonstrated in listing 39, key features of this algorithm include that users can configure arbitrary symbols to align by and lines can contain multiple vertically aligned columns. This algorithm

⁶ The actual implementation is 130 lines of code including helper methods, see <https://git.io/voIG4>

Algorithm 7: Vertical alignment, simplified algorithm

```
1  case class FormatLocation(formatToken: FormatToken, split: Split, state: State)
2  /** Returns true if location is eligible for vertical alignment */
3  def isCandidate(location: FormatLocation): Boolean
4  /** Returns true if all vertical alignment candidates in a and b match */
5  def allColumnsMatch(a: Array[FormatLocation], b: Array[FormatLocation]): Boolean
6  /** Returns map where the keys are (0 to block.length) and values are the
7   corresponding column index where all candidates should align */
8  def getMaxColumns(block: Vector[Array[FormatLocation]]): Map[Int, Int]
9
10 def getAlignTokens(
11   locations: Array[FormatLocation],
12   alignConfiguration: Map[String, Regex]): Map[Split, Int] = {
13   val finalResult = Map.newBuilder[Split, Int]
14   val lines: Array[Array[FormatLocation]] = getLines(locations)
15   var block = Vector.empty[Array[FormatLocation]]
16   for (formatLocations <- lines) {
17     val candidates: Array[FormatLocation] = formatLocations.filter(isCandidate)
18     if (block.isEmpty) { // Starting a new block.
19       if (candidates.nonEmpty) block = block :+ candidates
20     } else {
21       if (columnsMatch(block.last, candidates)) {
22         block = block :+ candidates
23       } else { // release alignment
24         val maxColumns = getMaxColumns(block)
25         for (line <- block) {
26           for ((tokenToAlign, columnIndex) <- line.zipWithIndex) {
27             finalResult += (tokenToAlign.split,
28                           maxColumns(columnIndex) - tokenToAlign.state.column)
29           }
30         }
31       }
32     }
33   }
34 }
```

has two main limitations. Firstly, all lines in a block must have an equal number of matching columns. Secondly, the algorithm can't infer vertical alignment forcing the user to explicitly configure which tokens should align.




Figure 4: Example heatmap with 5.121 visisted states

4 Tooling

This chapter describes the tools that we developed while designing and implementing algorithms for scalafmt. These tools were indispensable in giving us confidence that our algorithms worked as intended.

4.1 Heatmaps

Section 3.4 introduces several extensions to algorithm 1 that were required to get acceptable performance. In general, the extensions involved eliminating search states. To identify code patterns that triggered excessive search growth, we developed heatmaps.

Heatmaps are a visualization that displays which code regions are most frequently visited in the best-first search. Figure 4 shows an example heatmap. The intensity of the red color indicates how often a particular token was visited. A token highlighted by the lightest shade of red was visited twice while a token highlighted by the darkest shade of red was visited over 256 times. This figure demonstrates several of the optimizations discussed in section 3.4. Firstly, thanks to the `dequeueOnNewStatements` optimization, the background is plain white up to the `Seq`. The `Seq` gets visited twice, once when there’s a space after the `=` and once when there’s a newline. Secondly, due to the `OptimalToken` optimization, when the search gets into trouble it backtracks to the tuple `(0, 0)` instead of the `Seq[((Int, Int), Matri)]` type signature. Finally, because of the strategically placed comment at the end that exceeds the column limit, the search space grows out of bounds on the fourth argument triggering the `escapeInPathologicalCases` best-effort fallback. Without heatmaps, it would be a much greater challenge to get these insights.

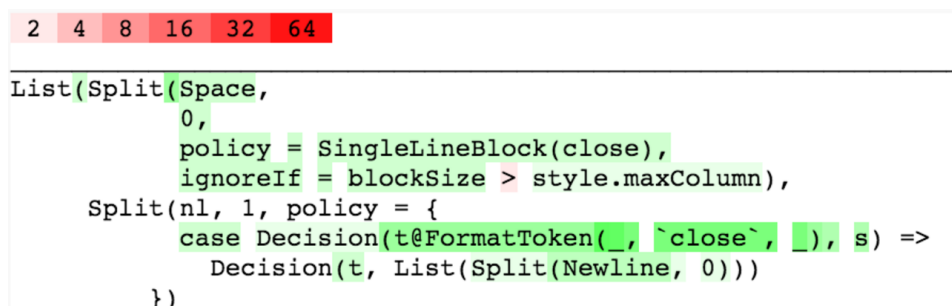


Figure 5: Example diff heatmap

However, these heatmaps gave us limited insights in how our changes affected the search space in the best-first search.

We developed an extension to heatmaps that allows us to visually compare the difference in search space between two versions of scalafmt. Figure 5 shows an example of such a diff report. The green background indicates that the new version of scalafmt makes fewer visits to those regions. Observe that the `>` operator has a background with a light shade of red. This means that the operator was visited more often in the new scalafmt version. A price well worth paying considering the overall shrink in search space. To produce diff heatmaps, we first persist to a database the statistics report needed to generate a single heatmap after each test run. Then, we generate the diff heatmap by fetching two reports and calculating the difference in visits per token. If the difference is negative for a particular token — meaning we visited said token fewer times — the background is highlighted green, otherwise red.

4.2 Property based testing

Property based tests verify that software implements a high-level specification. Unlike example based tests, the input is not manually curated by the software developer. Instead, the input is automatically generated. Property based tests played a vital role in development of scalafmt and gave us confidence that the algorithms in section 3 behave well in the real world. We collected a sample of 1.2 million lines of code from open source Scala projects available online. The sample was compressed into a 23mb zip

file⁷. Our test suite would download the sample and test again three properties: *can-format*, *AST integrity* and *idempotency*.

4.2.1 Can-format

The can-format property simply says that if the Scala compiler’s parser is able to parse the source input file, then scalafmt should be able to format the source file. Although this may seem like a trivial property, it was by far the most effective property at finding bugs in scalafmt. Most commonly, comments in the most unexpected placed would cause the best-first search to not reach the last token in the input. An overly strict Policy was usually the culprit of such bugs, which was easy to fix thanks to our tracing techniques described in section 3.3.1.

4.2.2 AST integrity

The AST integrity property says that the abstract syntax tree of the formatted source file should be identical to the abstract syntax tree of the original input. Recall from section 2.2 that scala.meta trees can be serialized into strings. We leverage this feature to test AST integrity. Algorithm 8 shows the code needed to test AST integrity. This property

Algorithm 8: AST integrity property

```
1 import scala.meta._
2 forAll { (code: String) =>
3   val beforeAST = code.parse[Source].show[Structure]
4   val afterAST = Scalafmt.format(code).parse[Source].show[Structure]
5   beforeAST == afterAST
6 }
```

caught several critical bugs. For example, in one case, scalafmt inserted a newline after the keyword **return**, breaking the semantics of the original source code. Moreover, this property highlighted the danger of enabling the **stripMargin** alignment. Since the **stripMargin** modified the contents of regular and interpolated string literals, the AST of the formatted output changed. A sign that **stripMargin** should be disabled by default. Finally,

⁷See <https://github.com/olafurpg/scalafmt/releases/download/v0.1.4/repos.tar.gz>

by adhering to AST integrity, we confine the scope of what scalafmt can do. This makes it easy for us to reject feature requests such as sorting imports. Such refactoring can be left to a separate tool, in our humble opinion.

4.2.3 Idempotency

The idempotency property says that if the output of formatting a source file twice should be identical to the output of formatting it once. This property is critical for scalafmt to be used as part of any continuous integration setup. It is not at all obvious that the algorithms in section 3 fulfill the idempotency property. Our experience reveals that it is in fact very easy to accidentally introduce non-idempotent formatting rules in the Router. We did not test for idempotency until the 0.2.3 release, after users reported non-idempotent formatting behavior in scalafmt. Yet, even after we started testing against idempotency in our comprehensive test-suite, we continued to receive issues with non-idempotent formatting. It turns out that 1.2 million lines of code is not a large enough sample to catch all property bugs.

5 Evaluation

Code formatting is inherently a subjective topic. This introduces a challenge when evaluating a code formatter. In this chapter, we will present measurements that we believe show the success of scalafmt. We do not measure how well software developers perceive scalafmt formatted code. Instead, we will focus on *performance benchmarks* and *user adoption*.

5.1 Performance benchmarks

In this chapter, we describe our test methodology, key metrics and the performance results.

5.1.1 Setup

The benchmarks are run on a Macbook Pro (Retina, 15-inch, Mid 2014) laptop with a quad-core 2.5 GHz Intel Core i7 processor, 256 KB L2 cache per core and 6 MB shared L3 cache. The laptop has 16 GB 1600 MHz DDR3 memory. The operating system is OS X El Capitan 10.11.5. We run the benchmarks from the scalafmt commit id [aff5e794](#) compiled against Scala 2.11.7, running on JVM version 8, update 91. For accurate measurements, all benchmarks are run with the OpenJDK Java Microbenchmark Harness (JMH)[[33](#)]. JMH takes into account a variety of parameters that affect performance on the JVM. The sbt-jmh[[26](#)] plugin makes it easy to integrate JMH with a Scala project.

5.1.2 Macro benchmark

The macro benchmarks is designed to get insight on how scalafmt performs in a continuous integration setup. For example, it is common to assert before code review that all source files are properly formatted. For this benchmark we format the entire Scala.js codebase. The codebase contains 915 source files and over 106 thousand lines of code, excluding blank lines and comments. For accurate measurements, we run five iterations of the macro benchmark. We compare the running time with Scalariform.

Table [1](#) shows the results from the macro benchmark. Scalafmt is almost 6x

Benchmark	Cores	Score	Error	Units
Parallel.scalafmt	4	14.616	± 0.632	s/op
Parallel.scalariform	4	2.810	± 0.641	s/op
Ratio		5.20		
Synchronous.scalafmt	1	35.654	± 0.459	s/op
Synchronous.scalariform	1	5.951	± 0.135	s/op
Ratio		5.99		

Table 1: Results from macro benchmark.

slower than Scalariform. Why is the performance gap so big? Is this gap acceptable for continuous integration setups?

We believe two factors contribute to the fact that scalafmt is 6x slower than Scalariform. Firstly, scalafmt’s formatting algorithm is more complex. Scalafmt may try thousands of different formatting layouts to find an optimal formatting output. In contrast, Scalariform’s formatting algorithm is linear. Secondly, scalafmt’s performance is tied with scala.meta, which is still a young library with opportunities to become better optimized. For example, profiling scalafmt reveals that 30% of the formatting time is dedicated to a pre-processing step — unrelated to the best-first search — that could be accomplished with minimal overhead during parsing inside scala.meta. In addition, preliminary results indicate that scalafmt may speed up yet another 30% by upgrading to the latest experimental release of scala.meta⁸.

We believe the current performance is acceptable for continuous integration setups. A typical diff in a code review touches only a few source files, and definitely far from 106 thousand lines of code as we test in this macro benchmark. Moreover, if the build server has multiple CPU cores the performance is less of an issue. However, we admit the current performance is far from excellent. We have reason to believe there are plenty of opportunities to improve scalafmt’s performance.

⁸ scala.meta has gone through several difficult non-source-compatible upgrades during the development of scalafmt.

25th	Median	Mean	75th	90th	95th	99th	Max
16	46	106	113	248	400	945	11723

Table 2: Percentiles of lines of code per file in micro benchmark.

5.1.3 Micro benchmark

The micro benchmark is designed to get insight into how scalafmt performs in an interactive software developer workflow. For example, it is common to configure a text editor to format source code on every save. Before we run the benchmark, we must find out how many lines of code a typical source file contains.

We performed a small study to learn the size of a typical source file. We collected a sample of 3.2 million lines of code from 33 open source Scala projects. Table 2 shows the distribution of file sizes in our sample. Observe that over 90 percent of all files are rather small, or under 250 lines of code. Only one percent of files contain more than 1.000 lines of code. Still, we assume developers spend quite a lot of time editing such large files.

Using the results from our small study, we choose to run the micro benchmark on four files of varying sizes: small (~ 50 LOC), medium (~ 300 LOC), large (~ 1.000 LOC) and extra large (~ 4.500 LOC). To minimize error margins, we run 10 warmup iterations followed by 10 measured iterations. As in the macro benchmark, we compare the running time with Scalariform. The micro benchmark is single threaded.

Table 3 shows the results from the micro benchmark. No surprise, scalafmt is again slower than Scalariform. Is this performance gap acceptable for interactive software development?

We believe this performance is usable for occasional code formatting, but not suitable for a workflow that formats on every save/compilation. Amazon famously showed that sales decreased by 1 percent for every 100ms increase in page load time[25]. We believe similar principles apply to scalafmt, every additional millisecond decreases the utility of scalafmt. Still, we believe with scalafmt’s formatting output is appealing enough to outweigh the slow performance. As we’ll discuss in the following section, our users seem to agree.

Benchmark	Score	Error	Units
ExtraLarge.scalafmt	1423.140	± 103.360	ms/op
ExtraLarge.scalariform	219.820	± 14.450	ms/op
Ratio	6.50		
Large.scalafmt	355.819	± 17.385	ms/op
Large.scalariform	39.324	± 3.395	ms/op
Ratio	9.05		
Medium.scalafmt	79.616	± 2.013	ms/op
Medium.scalariform	15.934	± 0.441	ms/op
Ratio	5.00		
Small.scalafmt	6.968	± 0.104	ms/op
Small.scalariform	1.176	± 0.025	ms/op
Ratio	5.93		

Table 3: Results from micro benchmark.

5.2 Adoption

Scalafmt has received quite some attention since its release in early March, three months ago. In this section we present the statistics we believe demonstrate that scalafmt is — despite its young age — already proving itself useful for the Scala community. All data points are as of June 9th, 2016.

5.2.1 Installations

Scalafmt has been installed over 6.500 times. Table 4 shows the installation numbers for each official distribution channel. IntelliJ is the JetBrains plugin repository⁹. Data is not available for how many users built scalafmt from source. The numbers represent absolute download numbers, not unique users. Observe that v0.2.5 was released 22 days ago, meaning it has been installed 80 times on average per day since its release. Considering the v0.2.5 installations numbers, it is fair to estimate that scalafmt currently has at least 1.000 users.

⁹ See <https://plugins.jetbrains.com/plugin/8236?pr=>




Figure 6: Installations by month by channel.

Table 4: Download numbers for scalafmt.

Channel	Version	Installations
IntelliJ	v0.2.5	847
	All	3.273
Maven	v0.2.5	788
	All	2.657
Github	v0.2.5	102
	All	929
Sum	v0.2.5	1.737
	All	6.859

Figure 6 shows the growth in installations by month. Observe that growth has doubled with each month. Github represented proportionally many installs in the first month but only represents a small fraction of installation in May. Maven installations quadrupled in May, taking the lead from the IntelliJ plugin in April.

5.2.2 Other

We present interesting data points from a variety of disparate data sources:

- The scalafmt code repository has received contributions from 8 external contributors. Several of these contributions added non-trivial features to scalafmt, including new configuration flags and extensions to the Router.
- 34 unique users, excluding the author, have opened a total of 138 tickets on the scalafmt issue tracker.
- The scalafmt Gitter¹⁰ instant messaging channel has 47 members. The channel is used to informally discuss bugs, new features and more.
- The user documentation website¹¹ has been visited 5.422 times with an average visit duration of 98 seconds. That makes 138 hours of

¹⁰See <https://gitter.im/olafurpg/scalafmt>

¹¹See <http://scalafmt.org>

combined visit time.

- Several popular open-source libraries have reformatted their codebases with scalafmt:
 - scala-native: <https://github.com/scala-native/scala-native>,
 - fetch: <https://github.com/47deg/fetch>,
 - psp-std: <https://github.com/paulp/psp-std>,
 - scala-js-dom: <https://github.com/scala-js/scala-js-dom>.

6 Discussion

6.1 Future work

6.2 Conclusion

References

- [1] *Akka*. URL: <http://akka.io/> (visited on 05/29/2016).
- [2] *Apache Spark™ - Lightning-Fast Cluster Computing*. URL: <http://spark.apache.org/> (visited on 05/29/2016).
- [3] *Bill Gosper*. URL: <http://gosper.org/bill.html> (visited on 05/31/2016).
- [4] Eugene Burmako. “Scala macros: let our powers combine!: on how rich syntax and static types work with metaprogramming”. In: *Proceedings of the 4th Workshop on Scala*. ACM. 2013, p. 3.
- [5] Eugene Burmako. *scala.meta*. <http://scalameta.org/>. (Accessed on 06/03/2016). Apr. 2016.
- [6] *CodeReviewComments for golang*. <https://github.com/golang/go/wiki/CodeReviewComments>. (Accessed on 06/01/2016). Oct. 2015.
- [7] Daniel Jasper. *clang-format - Automatic formatting for C++*. https://www.youtube.com/watch?v=s7JmdCfI_c. (Accessed on 06/02/2016).
- [8] David Copeland Daniel Spiewak. *Scala Style Guide*. <http://docs.scala-lang.org/style/>. (Accessed on 06/06/2016).
- [9] Edsger W. Dijkstra. “A note on two problems in connexion with graphs”. In: *Numerische mathematik* 1.1 (1959), pp. 269–271. URL: <http://www.springerlink.com/index/uu8608u0u27k7256.pdf> (visited on 06/01/2016).
- [10] Sébastien Doeraene. *Scala.js Coding Style*. 2015. URL: <https://github.com/scala-js/scala-js/blob/master/CODINGSTYLE.md> (visited on 05/28/2016).
- [11] Stuart Dreyfus. “Richard Bellman on the birth of dynamic programming”. In: *Operations Research* 50.1 (2002), pp. 48–51.
- [12] *ENSIME*. URL: <http://ensime.github.io/> (visited on 05/29/2016).
- [13] Ira Goldstein. *Pretty-printing Converting List to Linear Structure*. Massachusetts Institute of Technology. Artificial Intelligence Laboratory, 1973. URL: <http://www.softwarepreservation.net/projects/LISP/MIT/AIM-279-Goldstein-Pretty-Printing.pdf> (visited on 05/29/2016).

- [14] *Google C++ Style Guide*. URL:
<https://google.github.io/styleguide/cppguide.html> (visited on 05/28/2016).
- [15] Robert Griesemer. *gofmt - The Go Code Formatter*.
<https://golang.org/cmd/gofmt/>. (Accessed on 06/01/2016). June 2009.
- [16] Sam Halliday. *I don't have time to talk about formatting in code reviews. I want the machine to do it so I can focus on the design.* microblog. May 2016. URL:
<https://twitter.com/fommil/status/727879141673078785>
(visited on 05/29/2016).
- [17] Li Haoyi. *sourcecode - Scala library providing "source" metadata to your program*. <https://github.com/lihaoyi/sourcecode>. (Accessed on 06/04/2016). Feb. 2016.
- [18] R. W. Harris. “Keyboard standardization”. In: 10.1 (1956), p. 37.
URL: <http://massis.lcs.mit.edu/archives/technical/western-union-tech-review/10-1/p040.htm> (visited on 05/29/2016).
- [19] John Hughes. “The design of a pretty-printing library”. In: *Advanced Functional Programming*. Springer, 1995, pp. 53–96. URL:
http://link.springer.com/chapter/10.1007/3-540-59451-5_3
(visited on 01/06/2016).
- [20] Daniel Jasper. *clang-format*. Mar. 2014. URL:
<http://llvm.org/devmtg/2013-04/jasper-slides.pdf> (visited on 04/20/2016).
- [21] Daniel Jasper. *ClangFormat*. 2013. URL:
<http://clang.llvm.org/docs/ClangFormat.html> (visited on 06/01/2016).
- [22] Viktor Klang. *Code style should not be enforced by review, but by automate rewriting. Evolve the style using PRs against the rewriting config.* microblog. Feb. 2016. URL:
<https://twitter.com/viktorklang/status/696377925260677120>
(visited on 05/29/2016).
- [23] Jon Kleinberg and Éva Tardos. *Algorithm design*. Pearson Education India, 2006.

- [24] Donald E. Knuth and Michael F. Plass. “Breaking paragraphs into lines”. In: *Software: Practice and Experience* 11.11 (1981), pp. 1119–1184. URL: <http://onlinelibrary.wiley.com/doi/10.1002/spe.4380111102/abstract> (visited on 05/31/2016).
- [25] Ron Kohavi and Roger Longbotham. “Online experiments: Lessons learned”. In: *Computer* 40.9 (2007), pp. 103–105.
- [26] Konrad Malawski. *ktoso/sbt-jmh: "Trust no one, bench everything." - sbt plugin for JMH (Java Microbenchmark Harness)*. <https://github.com/ktoso/sbt-jmh>. (Accessed on 06/07/2016).
- [27] John McCarthy. “Recursive functions of symbolic expressions and their computation by machine, Part I”. In: *Communications of the ACM* 3.4 (1960), pp. 184–195. URL: <http://dl.acm.org/citation.cfm?id=367199> (visited on 05/31/2016).
- [28] Bob Nystrom. *dart_style - An opinionated formatter/linter for Dart code*. Sept. 2014. URL: https://github.com/dart-lang/dart_style (visited on 06/01/2016).
- [29] Bob Nystrom. *The Hardest Program I've Ever Written*. Sept. 2015. URL: <http://journal.stuffwithstuff.com/2015/09/08/the-hardest-program-ive-ever-written/> (visited on 04/14/2016).
- [30] Martin Odersky and Heather Miller. *The Scala Center*. Mar. 2016. URL: <http://www.scala-lang.org/blog/2016/03/14/announcing-the-scala-center.html> (visited on 05/29/2016).
- [31] Martin Odersky et al. *The Scala language specification*. 2004. URL: http://www-dev.scala-lang.org/old/sites/default/files/linuxsoft_archives/docu/files/ScalaReference.pdf (visited on 05/31/2015).
- [32] Bruno CdS Oliveira, Adriaan Moors, and Martin Odersky. “Type classes as objects and implicits”. In: *ACM Sigplan Notices*. Vol. 45. 10. ACM. 2010, pp. 341–360.
- [33] *OpenJDK: jmh*. <http://openjdk.java.net/projects/code-tools/jmh/>. (Accessed on 06/07/2016).

- [34] Dereck C. Oppen. “Prettyprinting”. In: *ACM Trans. Program. Lang. Syst.* 2.4 (Oct. 1980), pp. 465–483. ISSN: 0164-0925. DOI: [10.1145/357114.357115](https://doi.org/10.1145/357114.357115). URL: <http://doi.acm.org/10.1145/357114.357115> (visited on 04/18/2016).
- [35] Judea Pearl. “Heuristics: intelligent search strategies for computer problem solving”. In: (1984). URL: <http://www.osti.gov/scitech/biblio/5127296> (visited on 06/01/2016).
- [36] Tiark Ropf and Nada Amin. “From F to DOT: Type Soundness Proofs with Definitional Interpreters”. In: *arXiv:1510.05216 [cs]* (Oct. 2015). arXiv: 1510.05216. URL: <http://arxiv.org/abs/1510.05216> (visited on 05/28/2016).
- [37] Matt Russell. *Scalariform*. 2010. URL: <http://scala-ide.org/scalariform/> (visited on 05/28/2016).
- [38] *sbt - The interactive build tool*. URL: <http://www.scala-sbt.org/> (visited on 05/28/2016).
- [39] *scala-native/scala-native*. URL: <https://github.com/scala-native/scala-native> (visited on 05/29/2016).
- [40] *Scala.js*. URL: <http://www.scala-js.org/> (visited on 05/29/2016).
- [41] R. S. Scowen et al. “SOAP—A program which documents and edits ALGOL 60 programs”. In: *The Computer Journal* 14.2 (1971), pp. 133–135. URL: <http://comjnl.oxfordjournals.org/content/14/2/133.short> (visited on 05/29/2016).
- [42] *Stack Overflow Developer Survey 2015*. URL: <http://stackoverflow.com/research/developer-survey-2015> (visited on 05/29/2016).
- [43] S. Doaitse Swierstra and Olaf Chitil. “Linear, bounded, functional pretty-printing”. In: *Journal of Functional Programming* 19.01 (Jan. 2009), pp. 1–16. ISSN: 1469-7653. DOI: [10.1017/S0956796808006990](https://doi.org/10.1017/S0956796808006990). URL: http://journals.cambridge.org/article_S0956796808006990 (visited on 04/20/2016).

- [44] Mark Van Den Brand and Eelco Visser. “Generation of formatters for context-free languages”. In: *ACM Transactions on Software Engineering and Methodology (TOSEM)* 5.1 (1996), pp. 1–41. URL: <http://dl.acm.org/citation.cfm?id=226156> (visited on 01/06/2016).
- [45] Philip Wadler. “A prettier printer”. In: *The Fun of Programming, Cornerstones of Computing* (2003), pp. 223–243. URL: <http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.19.635&rep=rep1&type=pdf> (visited on 04/20/2016).
- [46] Hyrum Wright et al. “Large-Scale Automated Refactoring Using ClangMR”. In: (2013). URL: <https://research.google.com/pubs/pub41342.html> (visited on 04/21/2016).
- [47] Reynold Xin. *Spark Scala Style Guide*. Mar. 2015. URL: <https://github.com/databricks/scala-style-guide> (visited on 06/01/2016).
- [48] Phillip M. Yelland. *A New Approach to Optimal Code Formatting*. Tech. rep. Google, inc., 2016. URL: <http://research.google.com/pubs/archive/44667.pdf> (visited on 04/20/2016).